

Kubernetes i CloudNative introduction



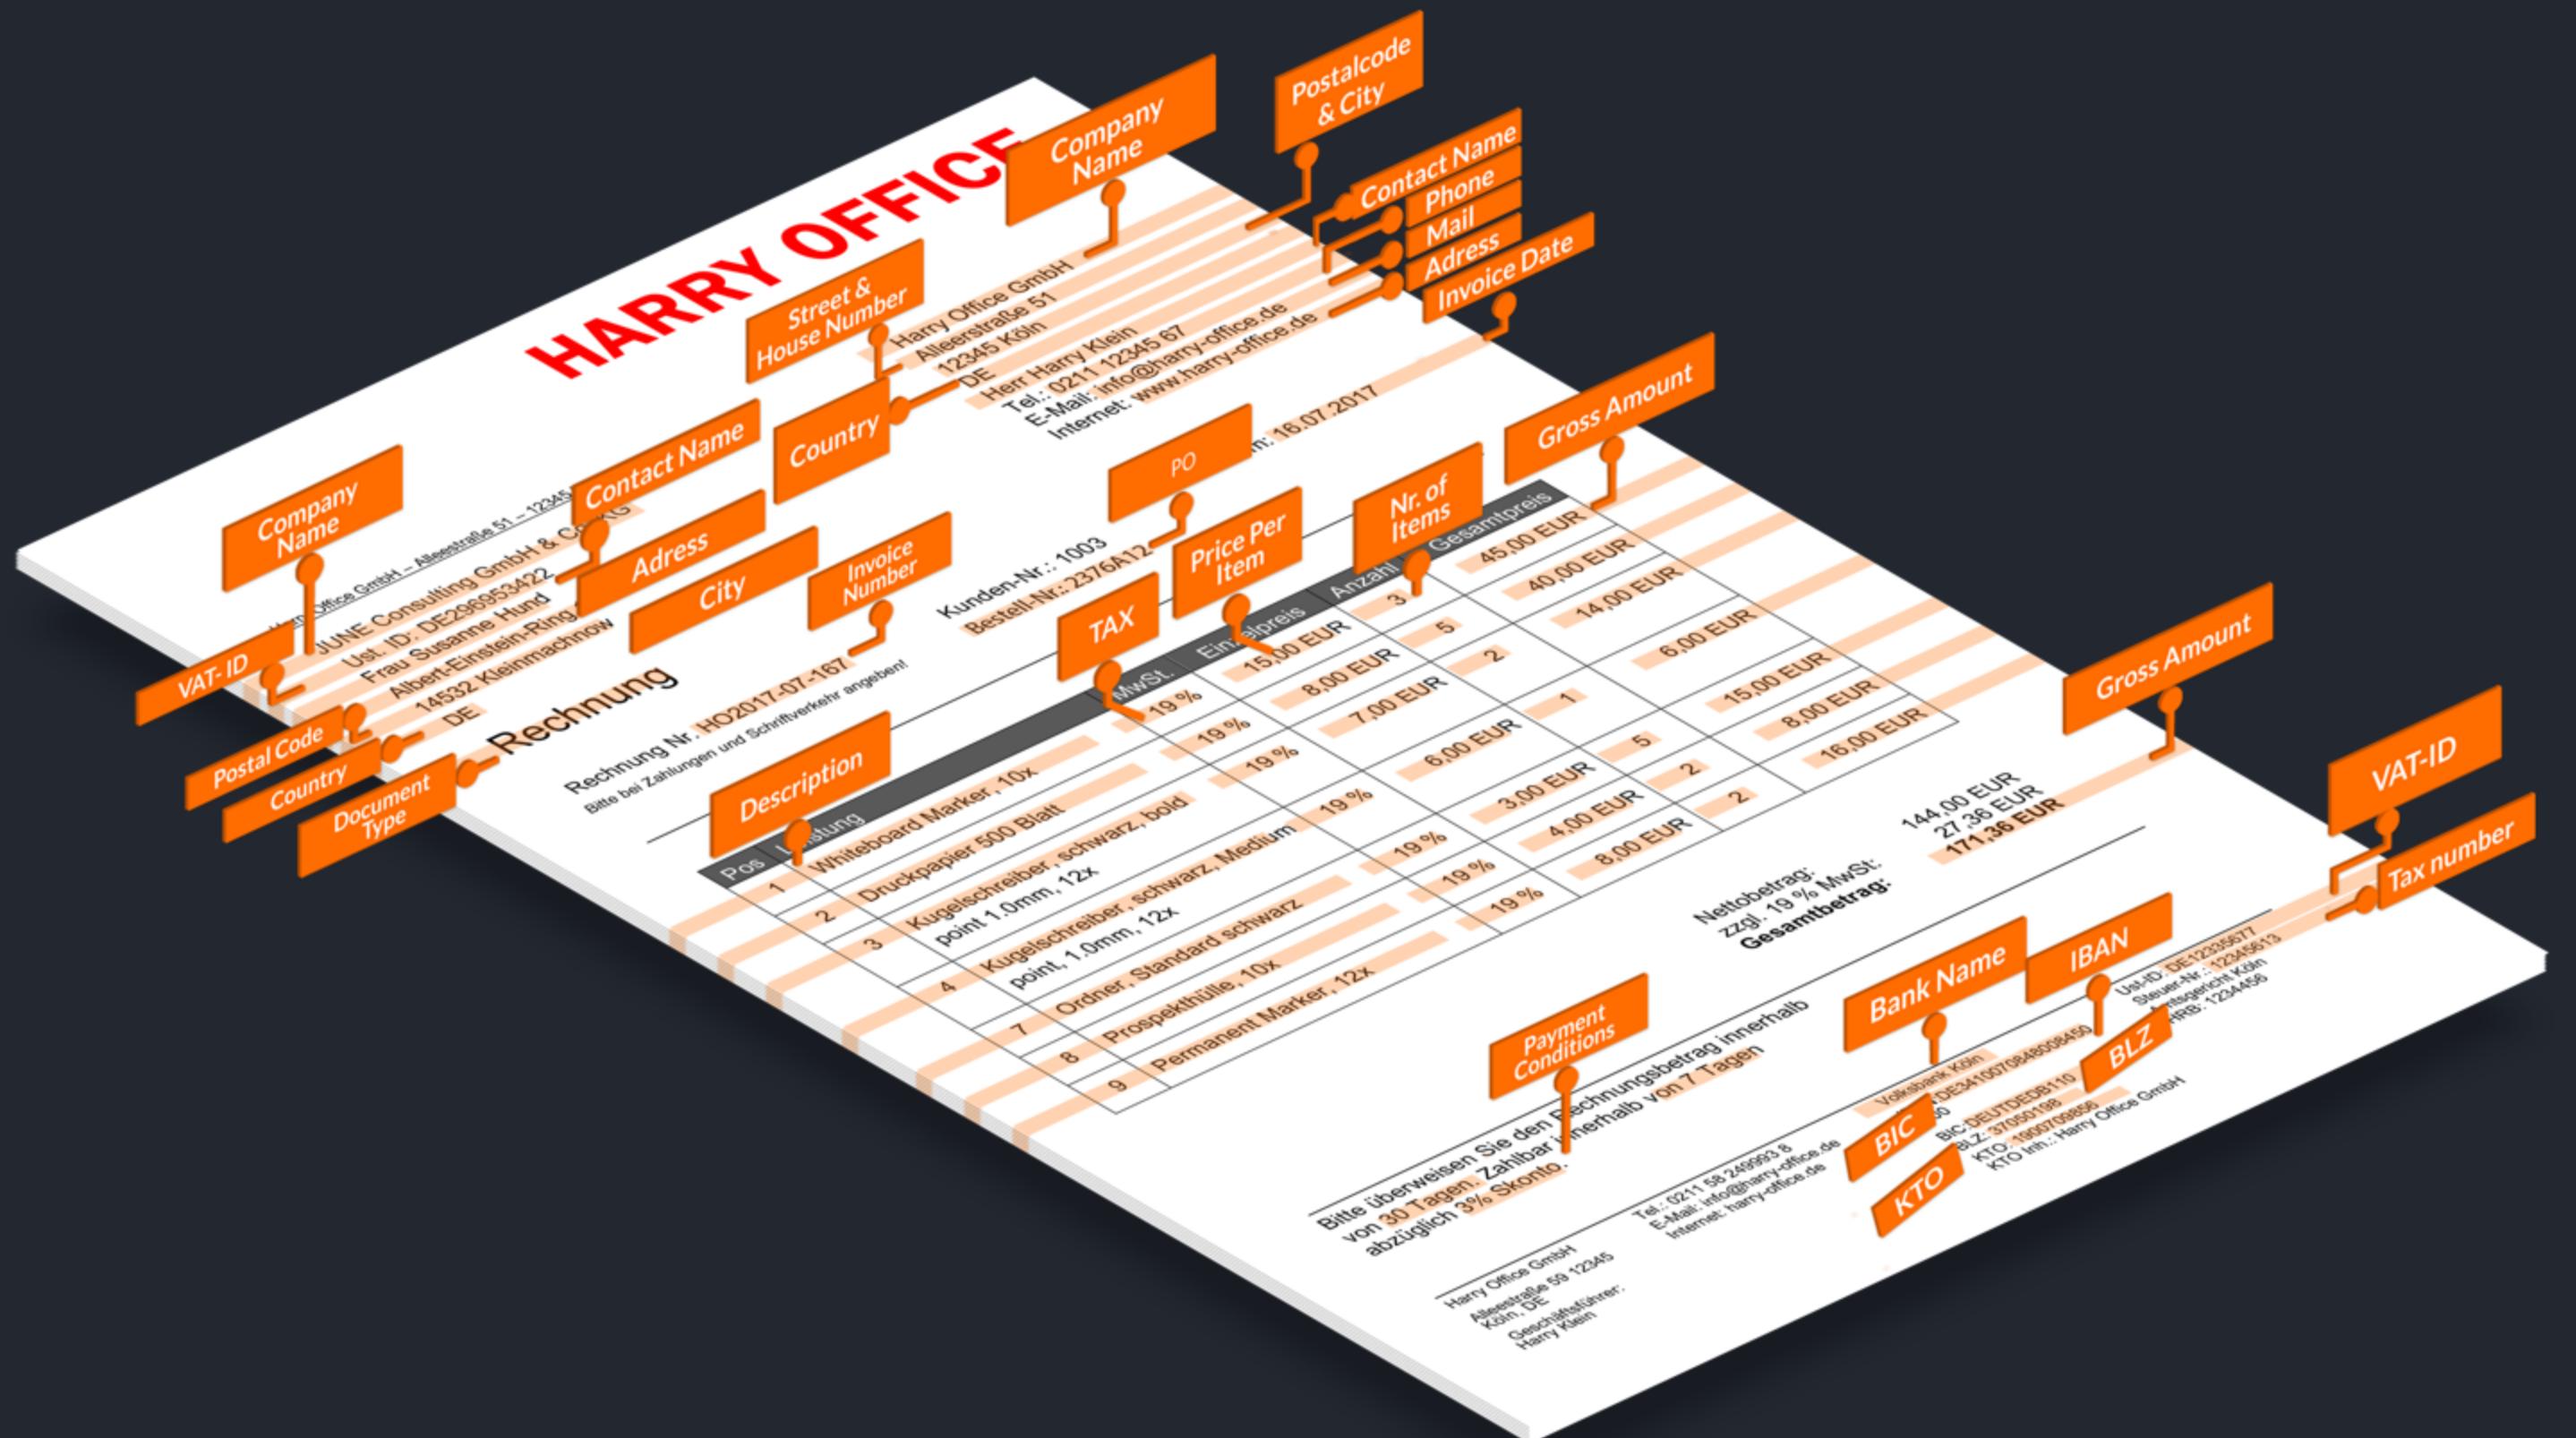
Wojciech Barczyński - [SMACC.io](#) | [Hypatos.ai](#)
27 March 2019

WOJCIECH BARCZYŃSKI

- Software Developer & System Engineer
- Now at Machine-Learning FinTech
- Since 2016 working with K8S

Github: [wojciech12](#) | Linked: [IN](#) | HP: [wbarczynski.pl](#) | T: [@wbarczynski](#)

Problem SMACC solves



Hypatos



SMACC

Beginnings



... - 2014

- Everybody talks *services* but....
- No simple platform for operating services
- Mesosphere!
- Docker swarm not perfect
- We learn so much in Google (Borg) and other companies

Docker like RoR change

New Linux, New Application Server!



OnPremise



CLOUD NATIVE COMPUTING FOUNDATION



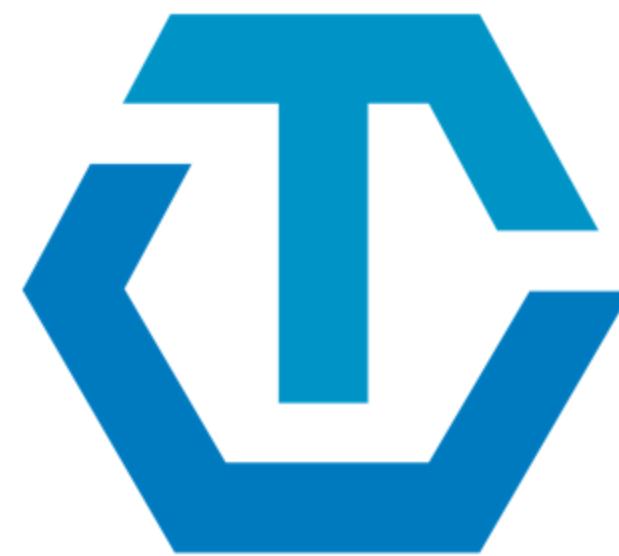
kubernetes



Prometheus



fluentd



OPENTRACING



Istio



<https://github.com/cncf/artwork>

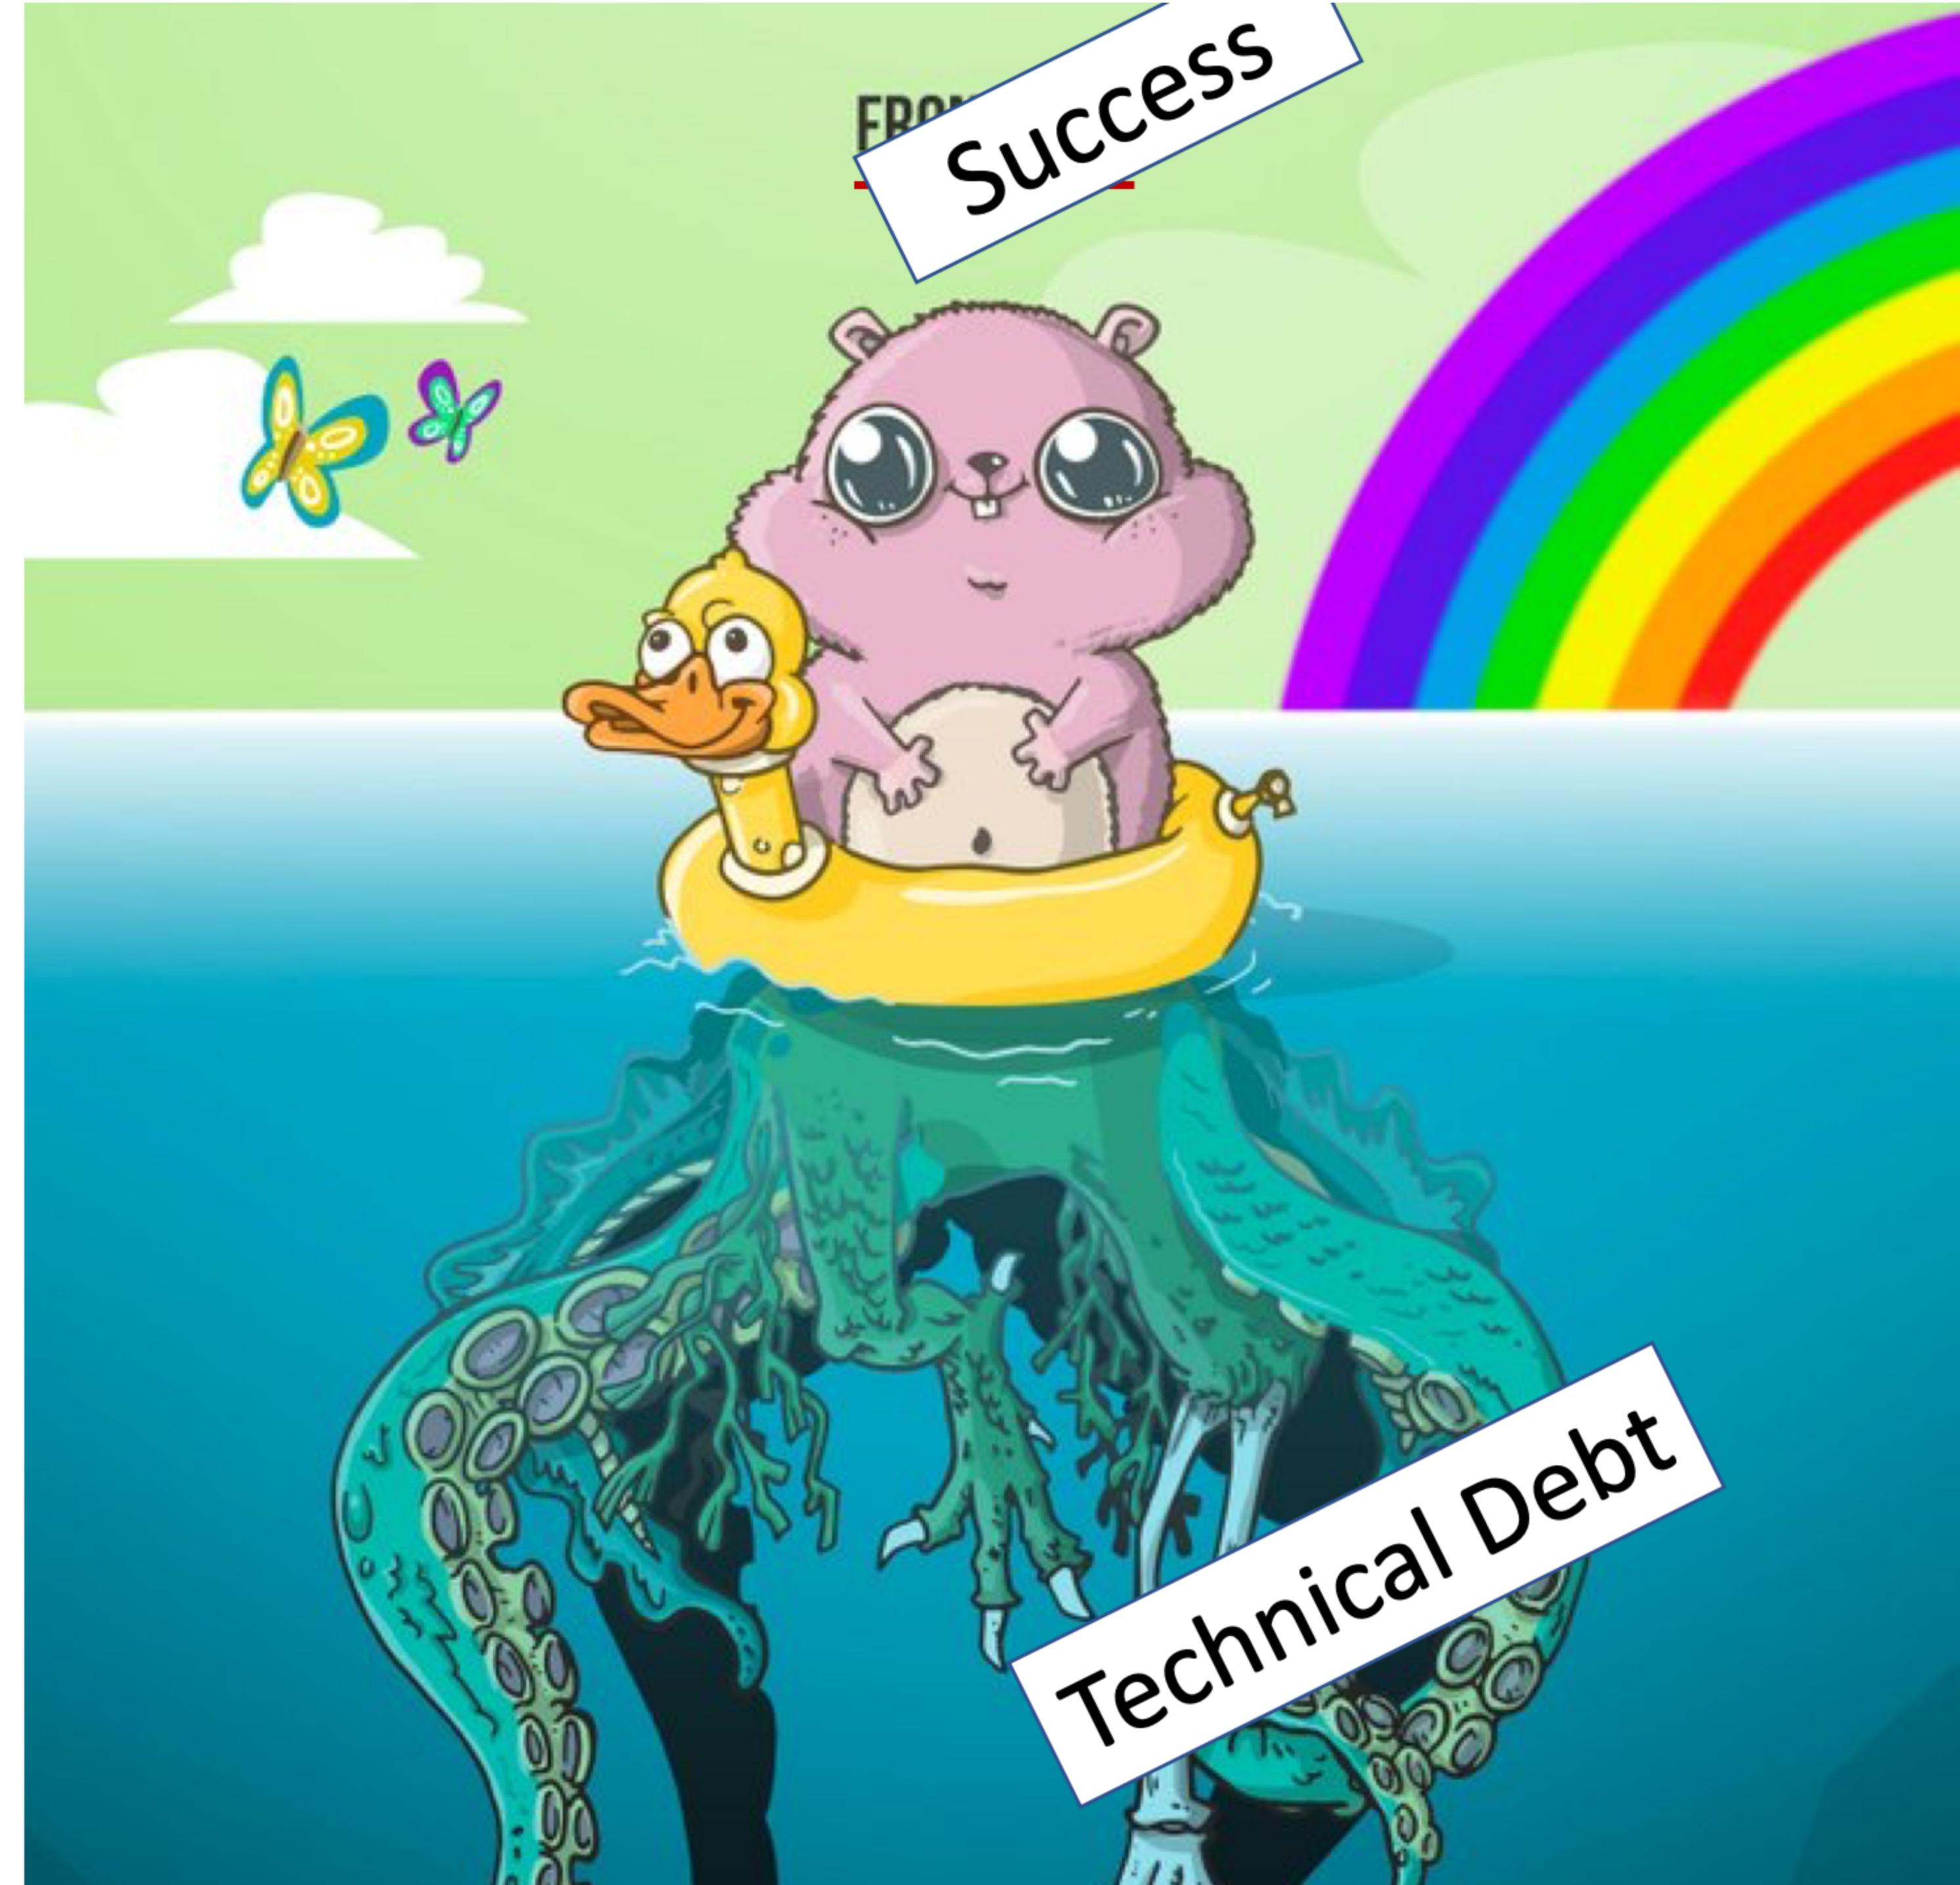
Slow delivery

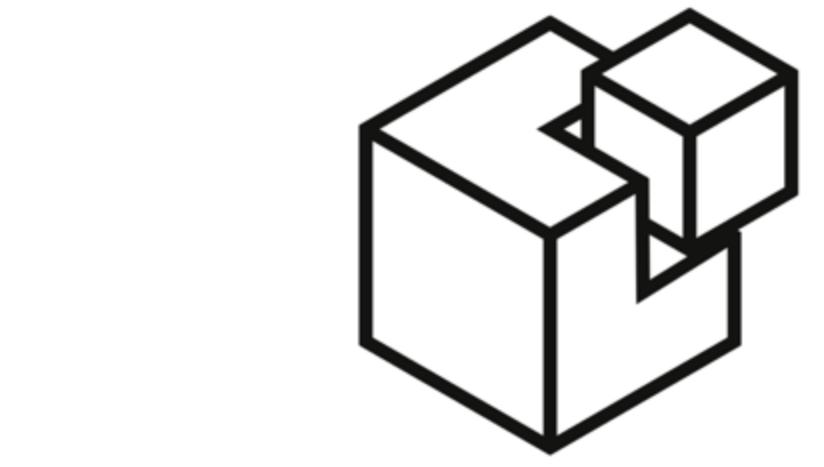
Continuous Deployment?

Fear

Frustration

XX% Idle Machines





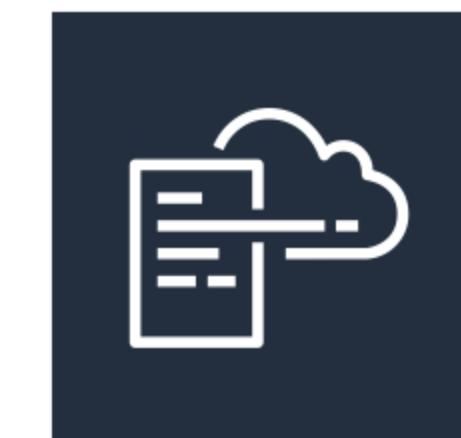
SALTSTACK



ANSIBLE



CHEF



AWS
CloudFormation

2016



HashiCorp
Terraform

Black (Blue) Box

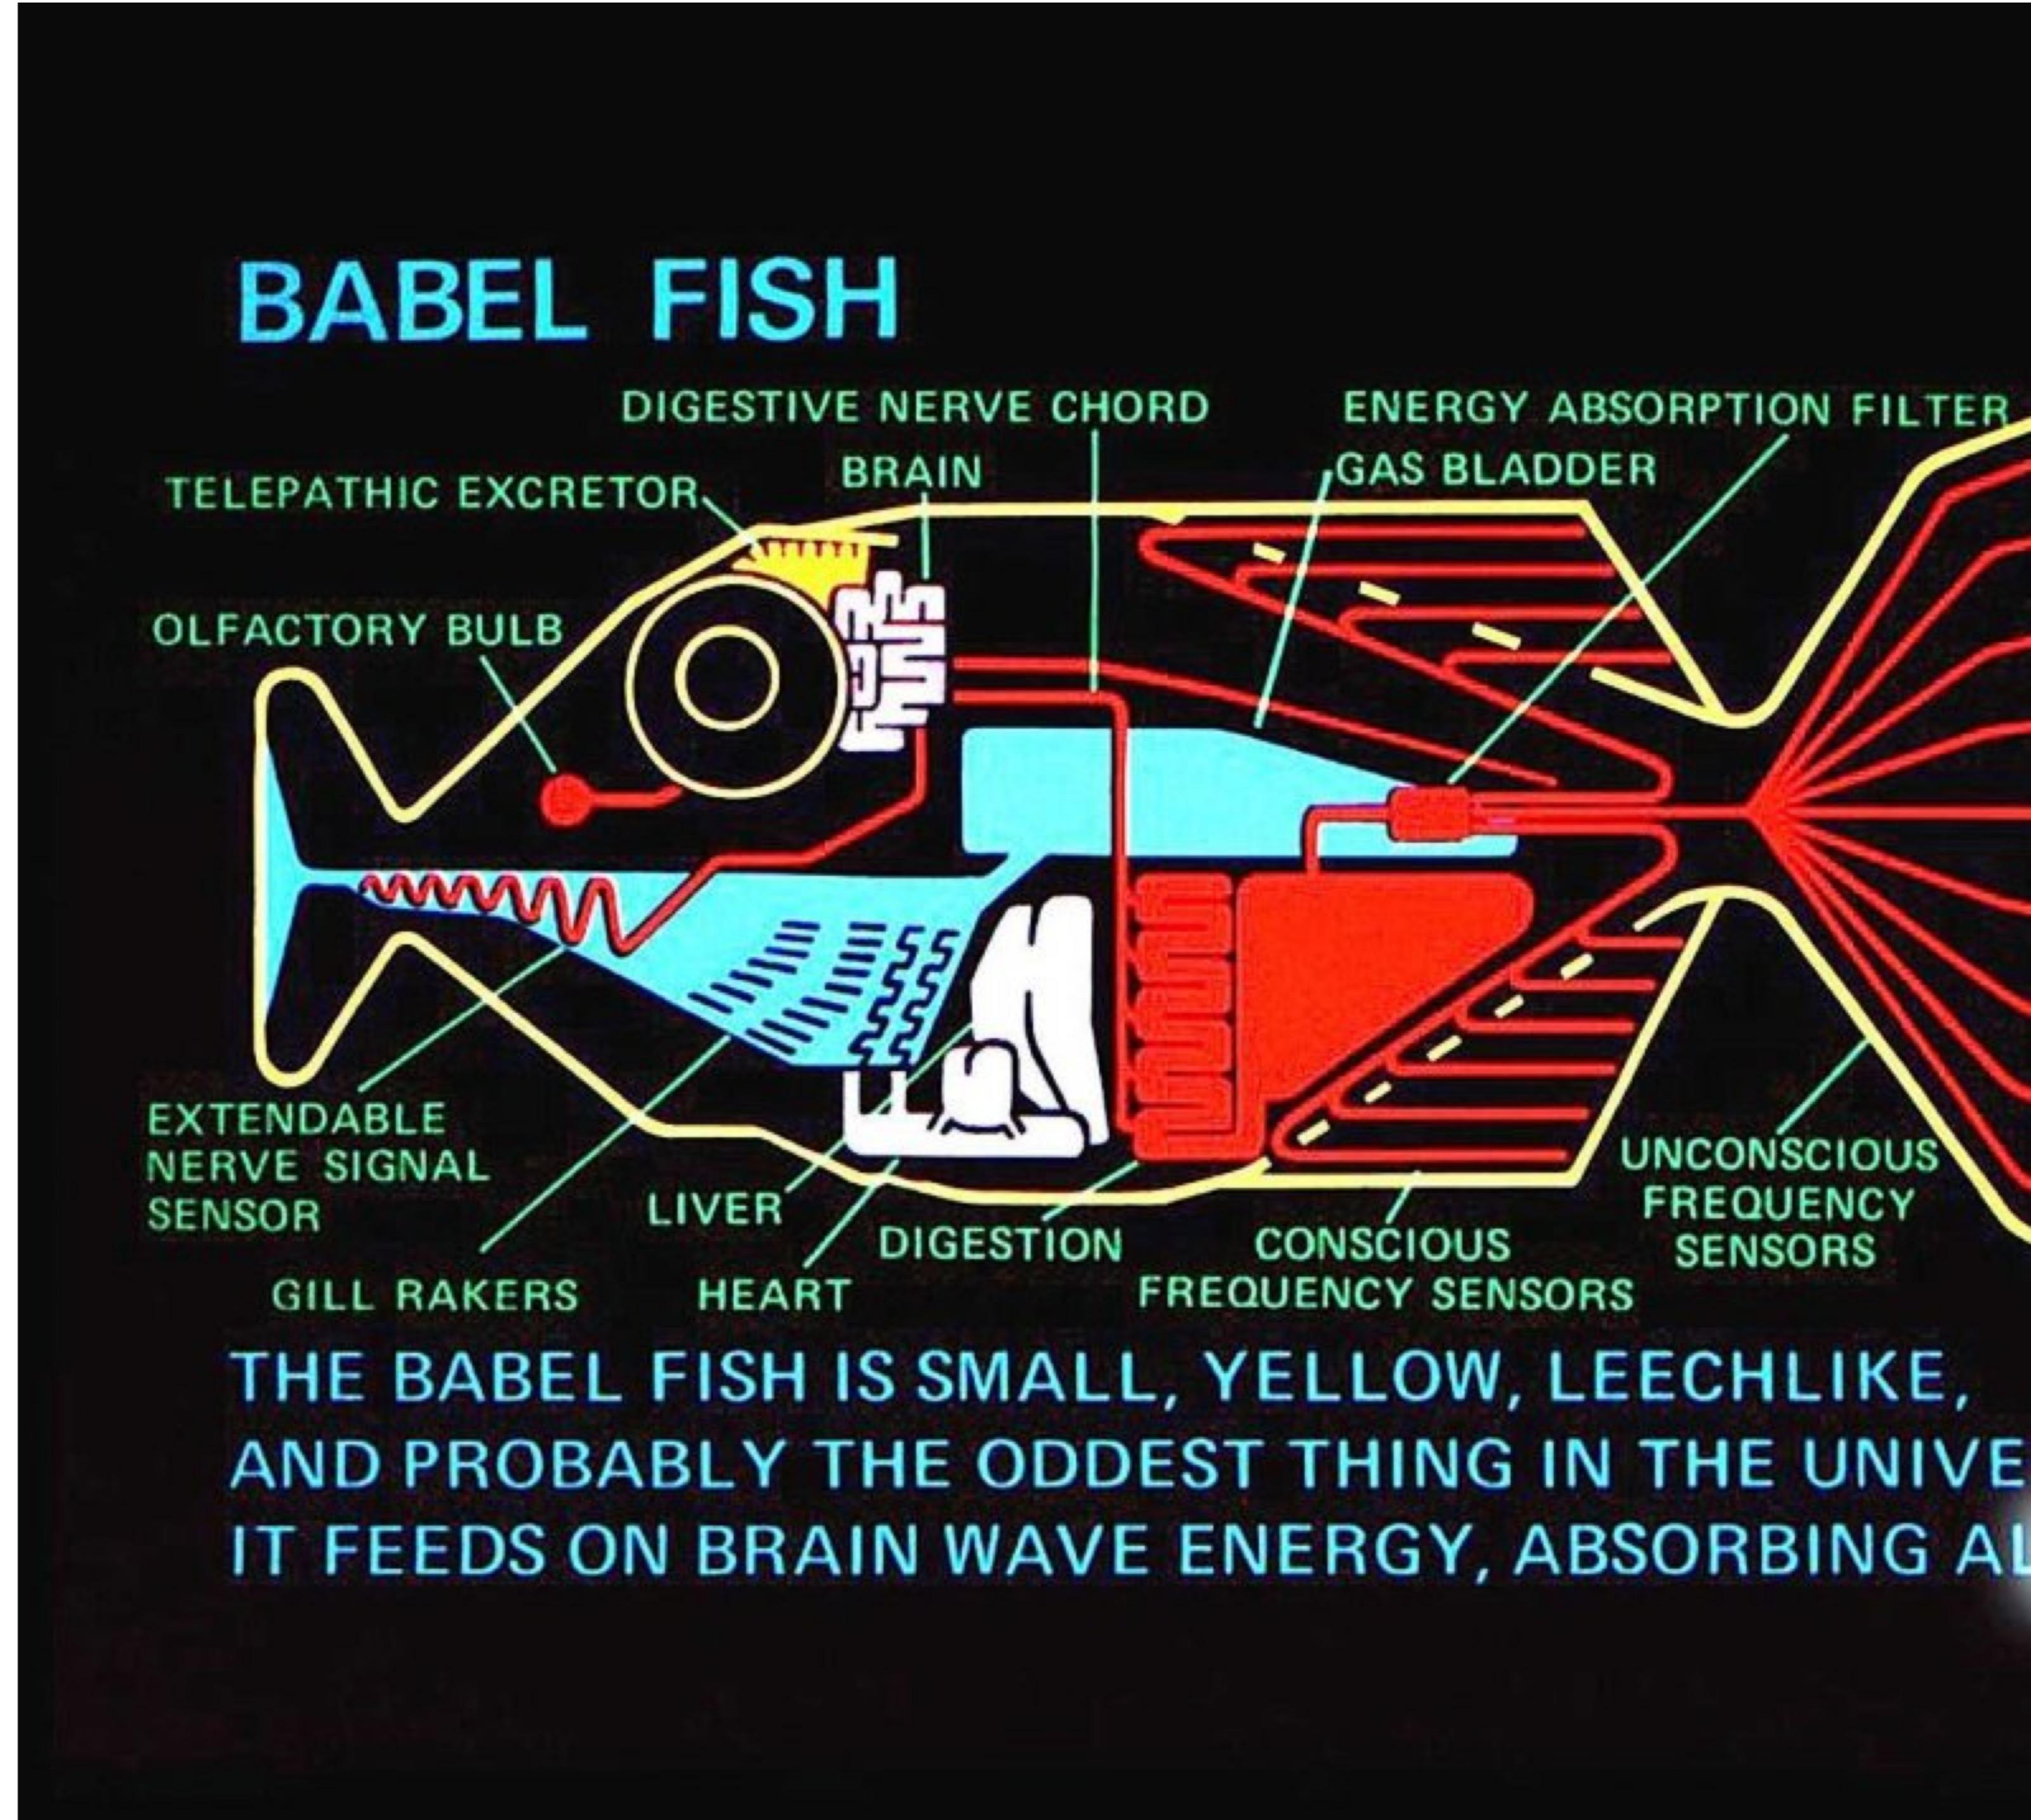
Infrastructure (almost) invisible

Easy* Continuous Deployment



[https://en.wikipedia.org/wiki/File:Dr_Who_\(316350537\).jpg](https://en.wikipedia.org/wiki/File:Dr_Who_(316350537).jpg)

Common
Language
Artifacts
Platform

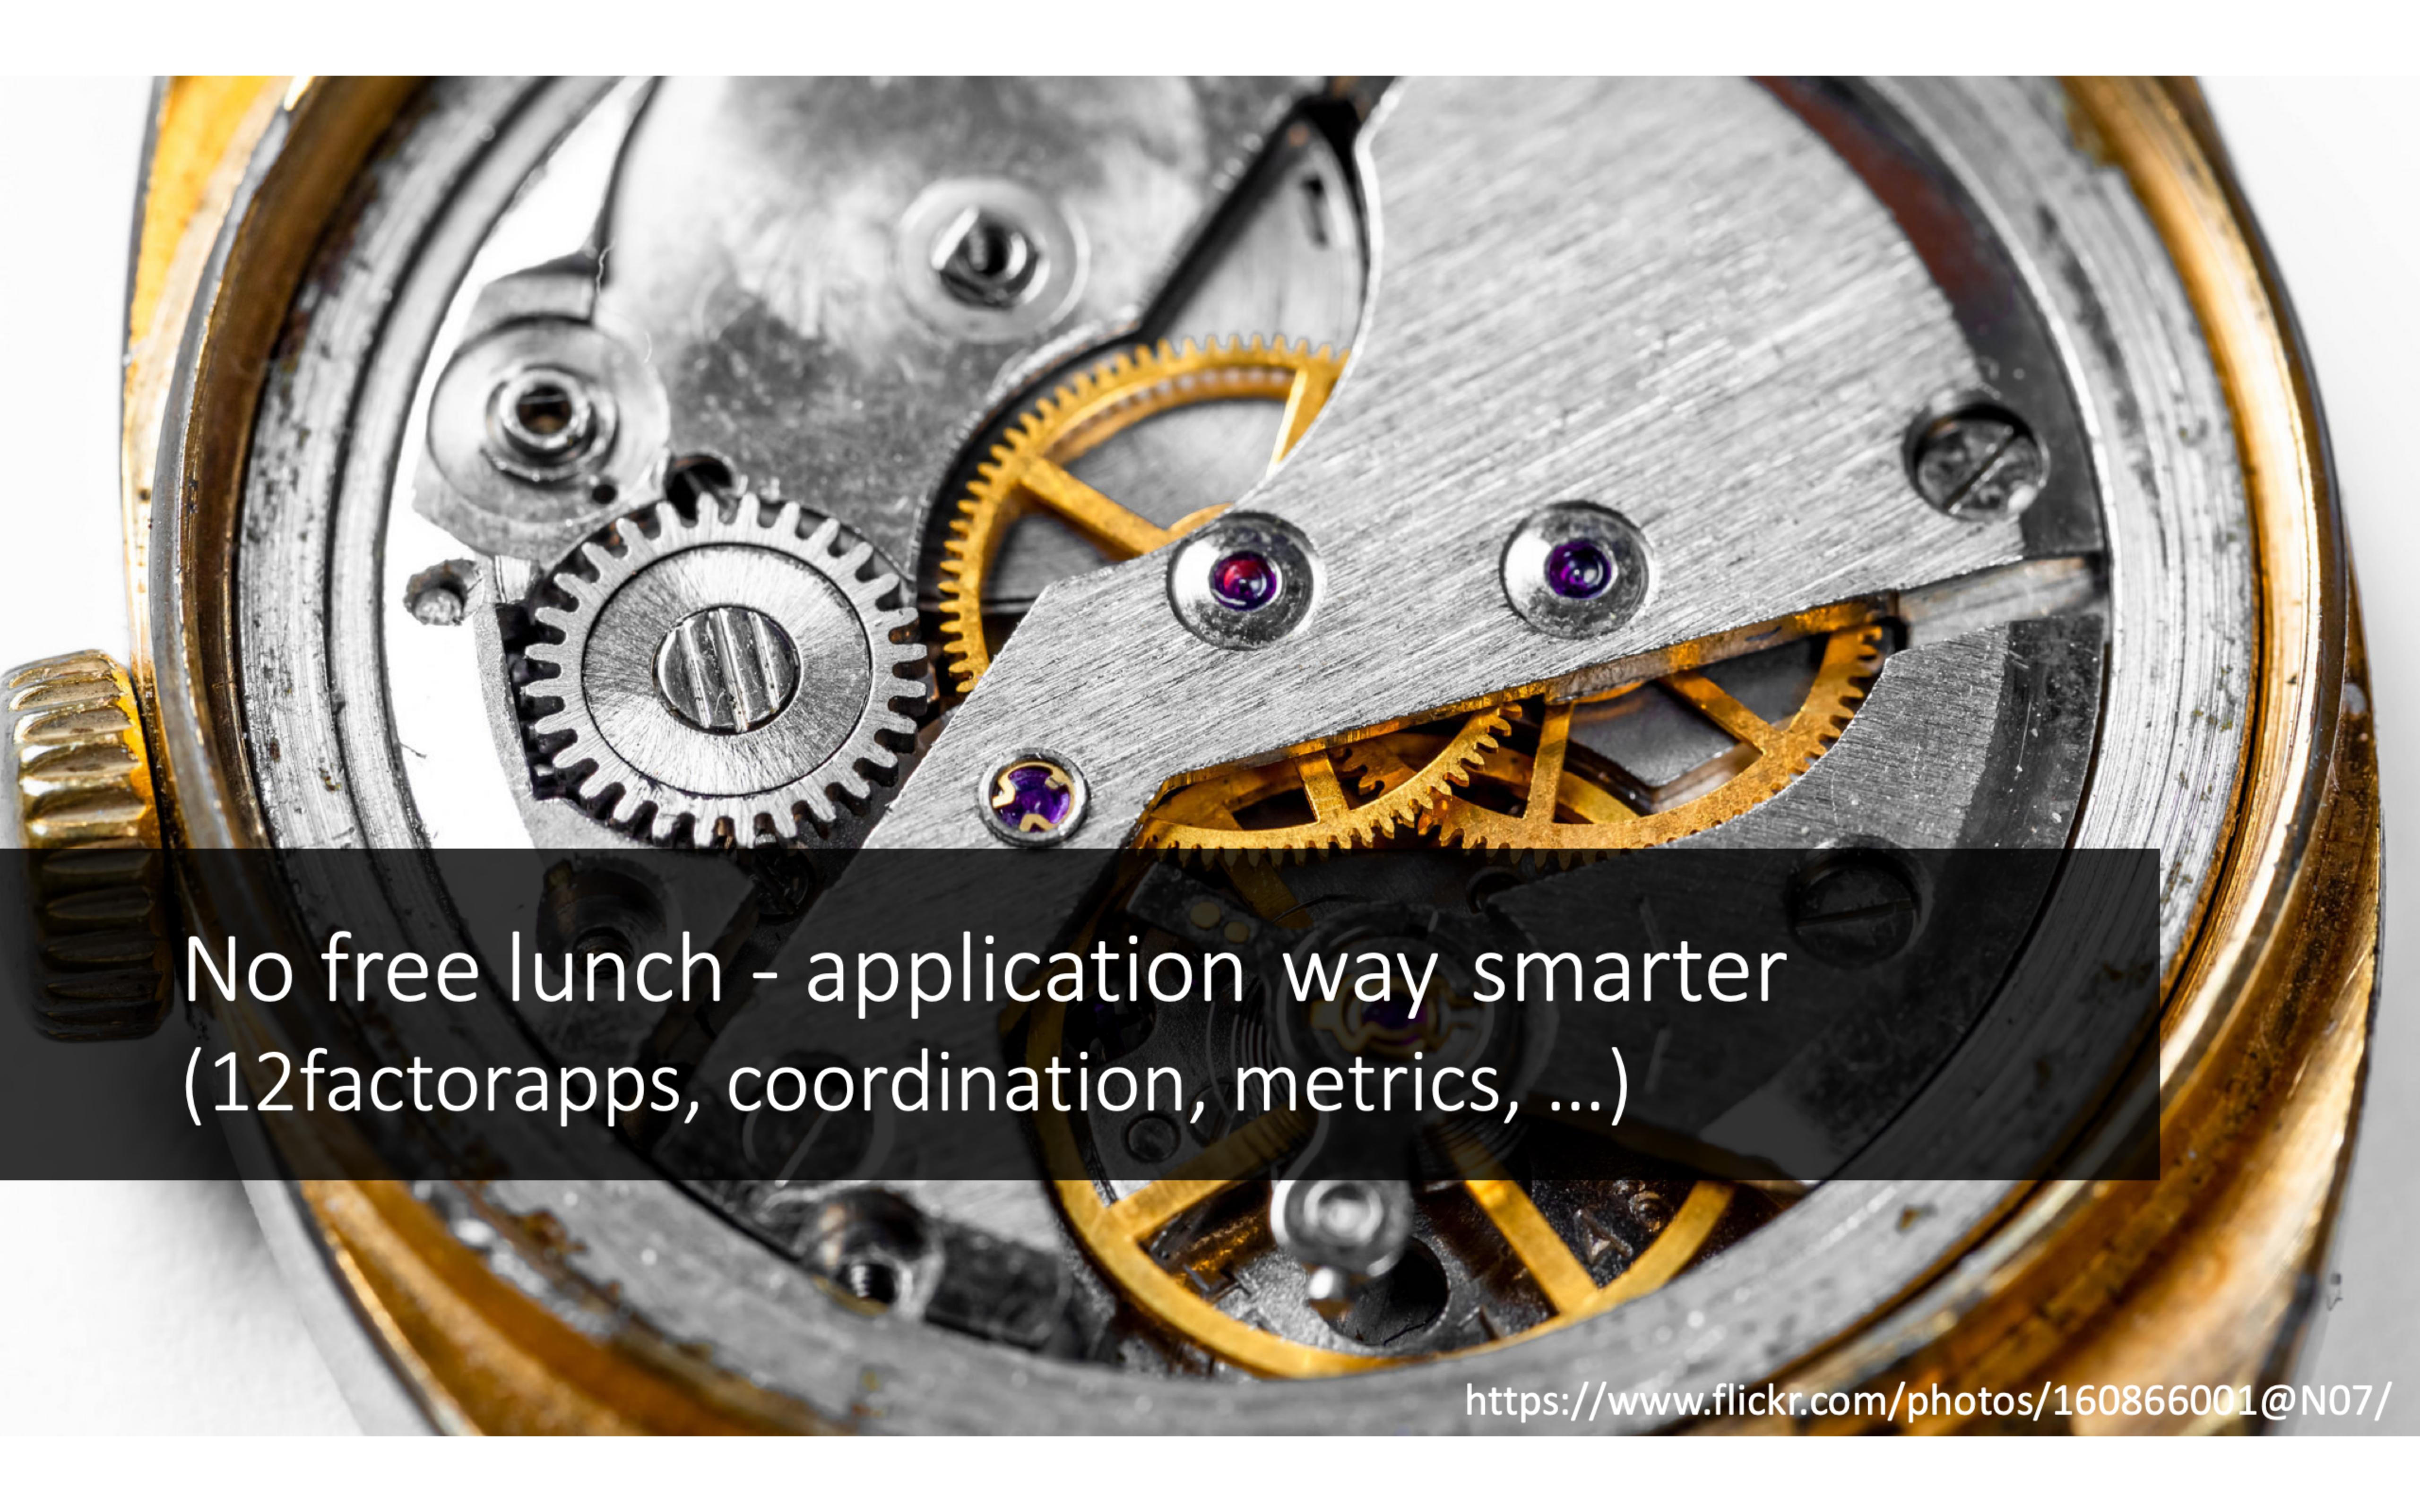


Learn-as-you-go

1. Deploy Cloud-Native app
2. Make a Hell of Mistakes
3. Get it right or Postpone



envoy



No free lunch - application way smarter
(12factorapps, coordination, metrics, ...)

KUBERNETES

- Service oriented
- Battery for 12factor apps
- Orchiestration
- Focus on simplicity [*]

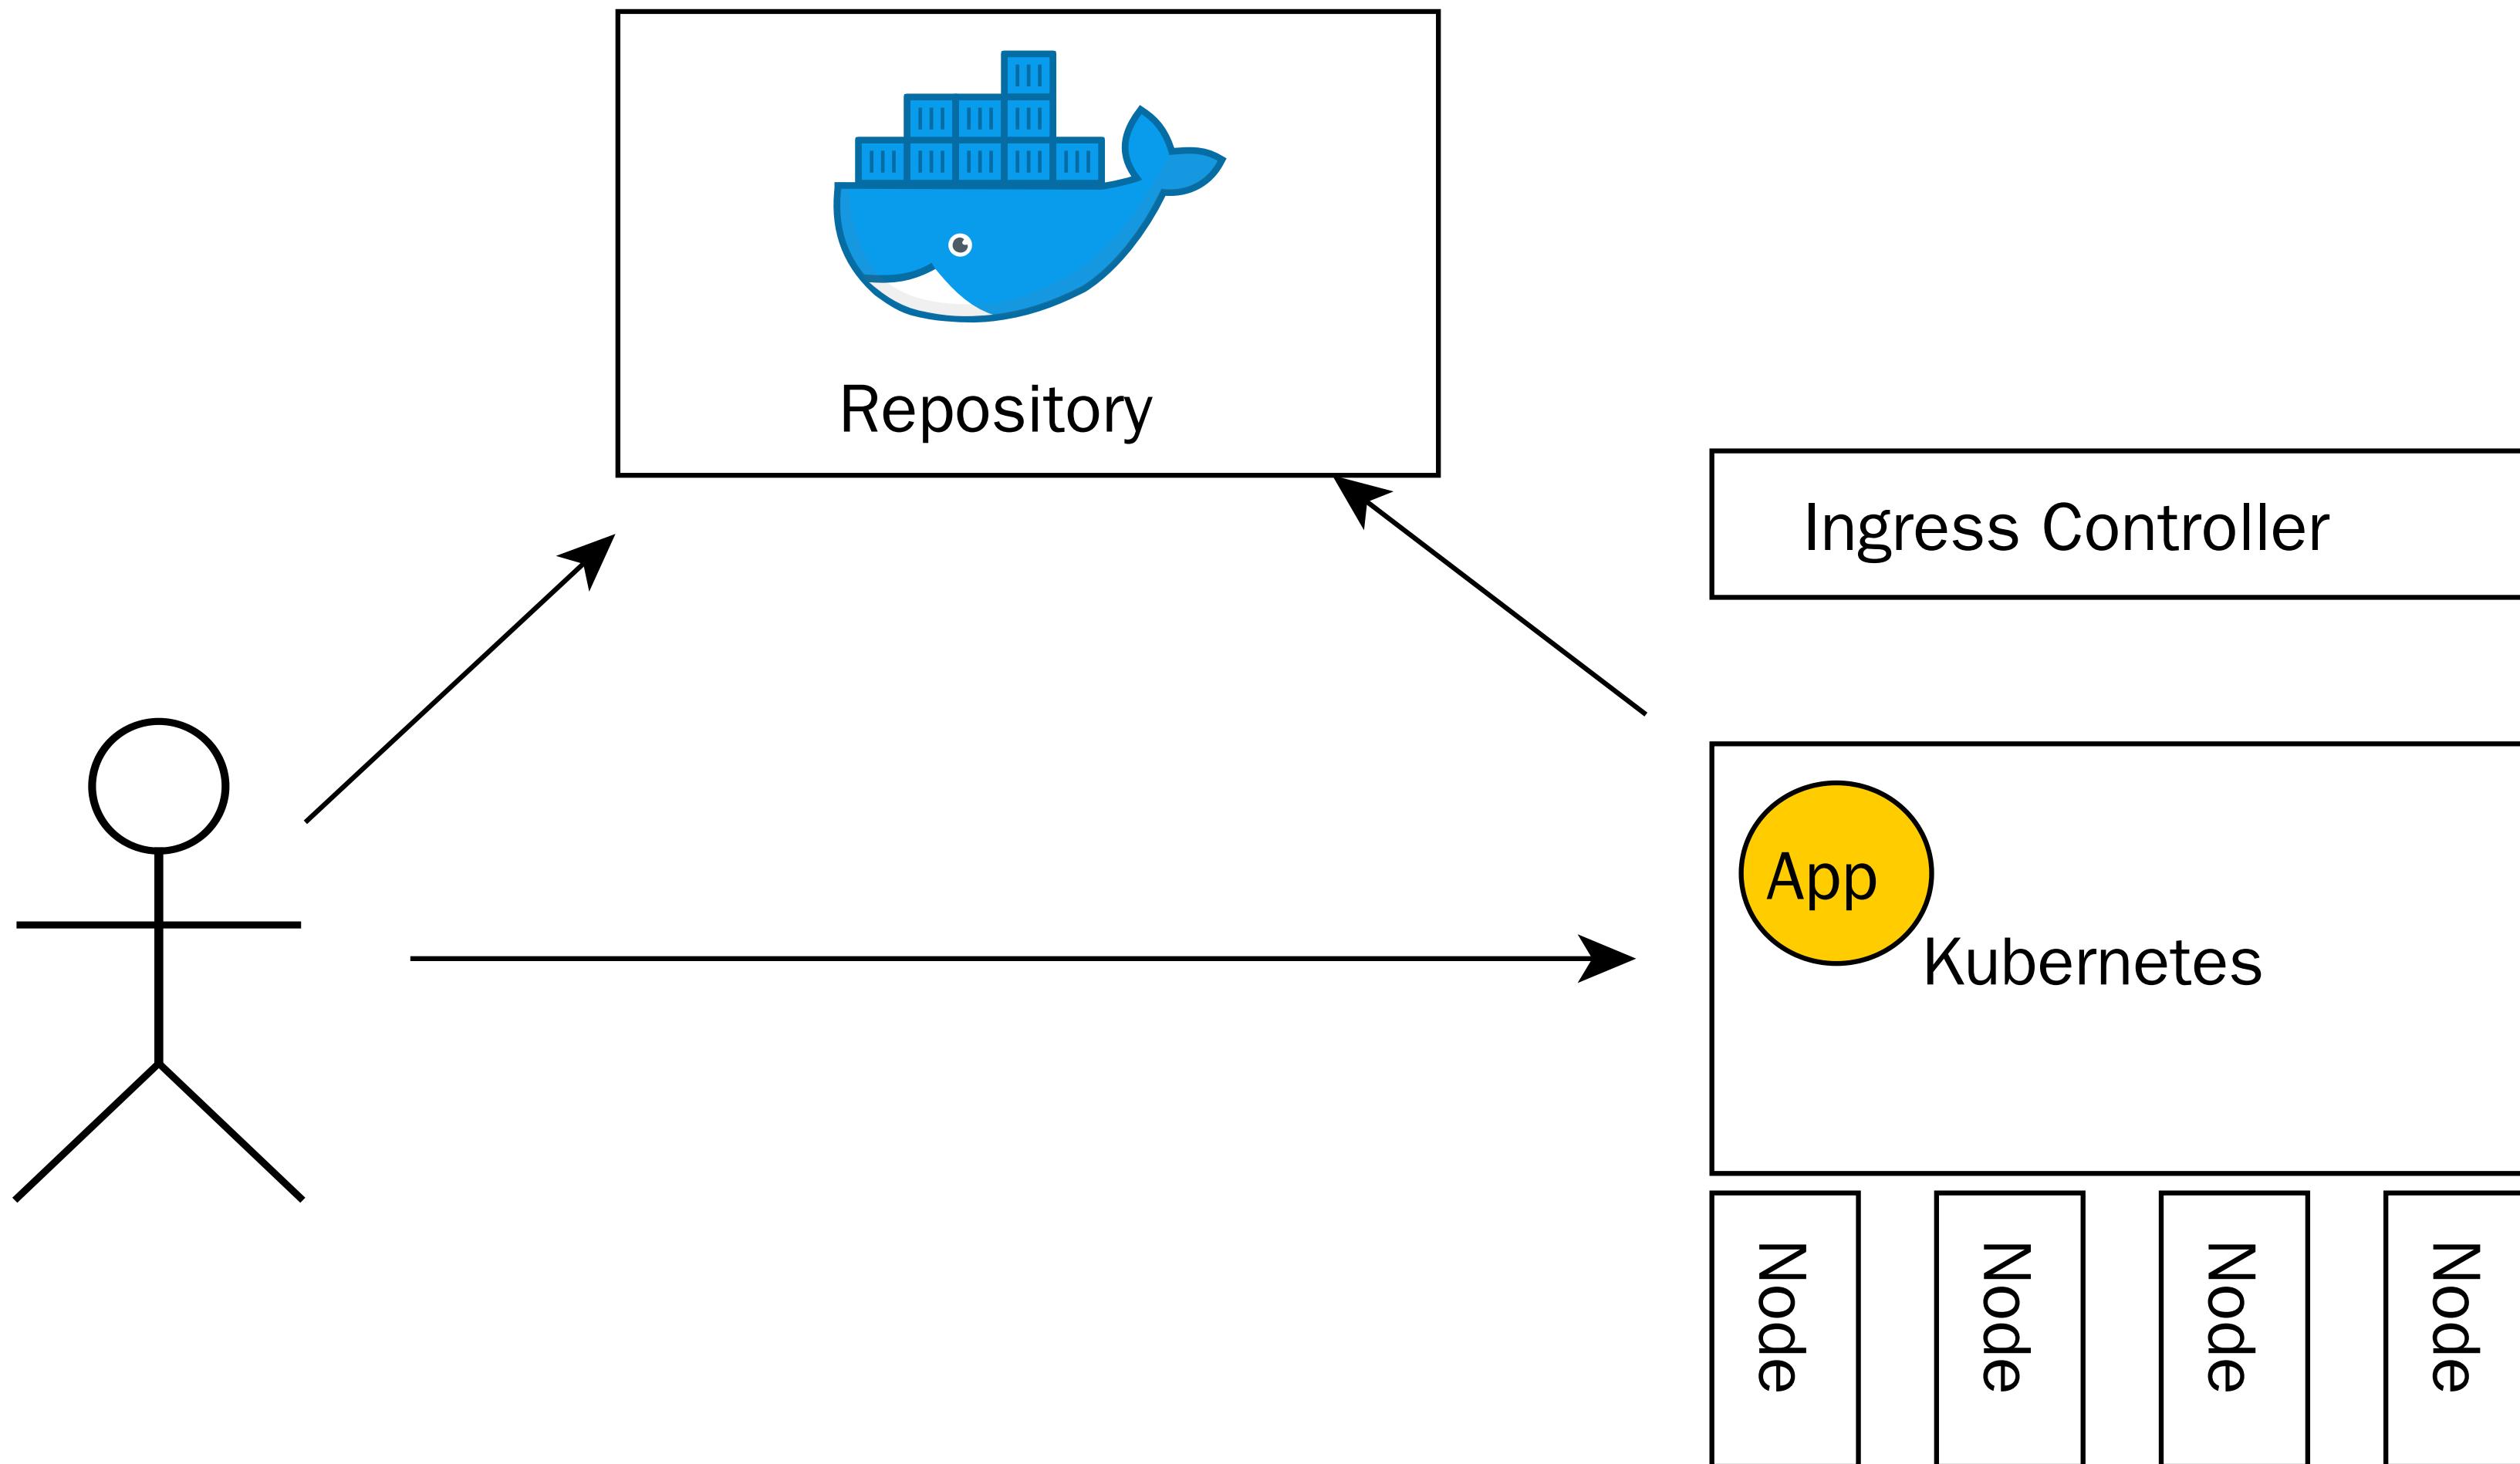
[*] In spirit what Docker brought us

The twelve factor apps

- Services to operate service, test, etc.
- Services with low TCO.
- Published by Heroku
- ps. No free lunch

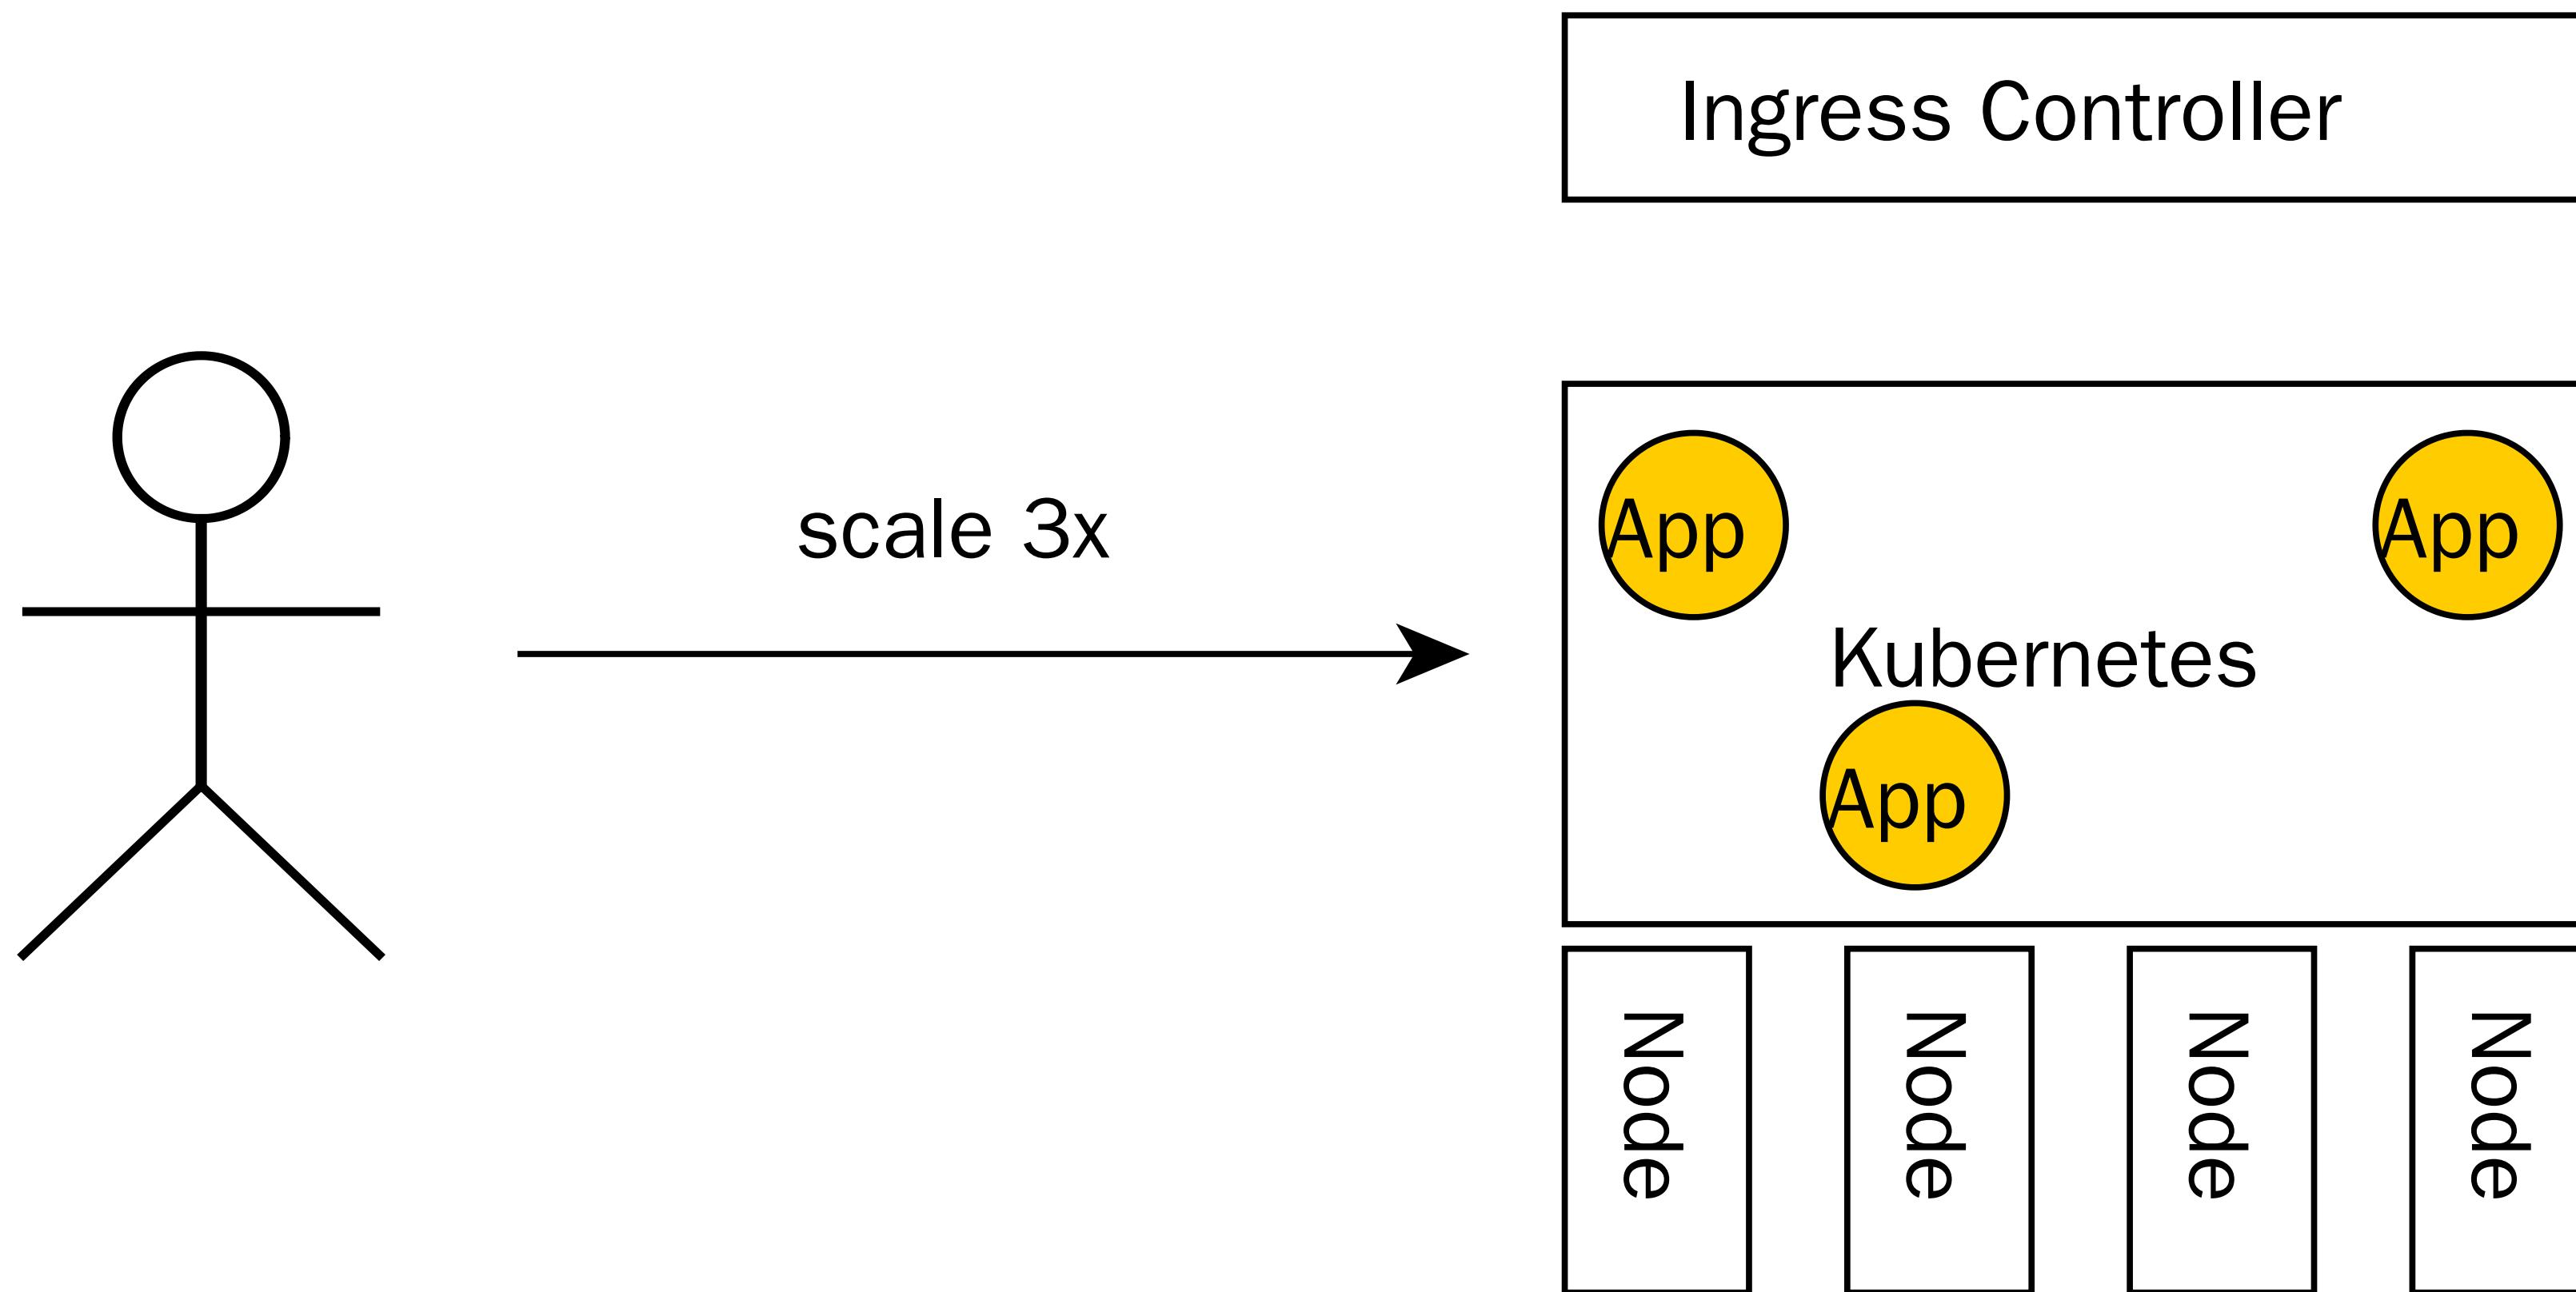
<https://12factor.net>

KUBERNETES



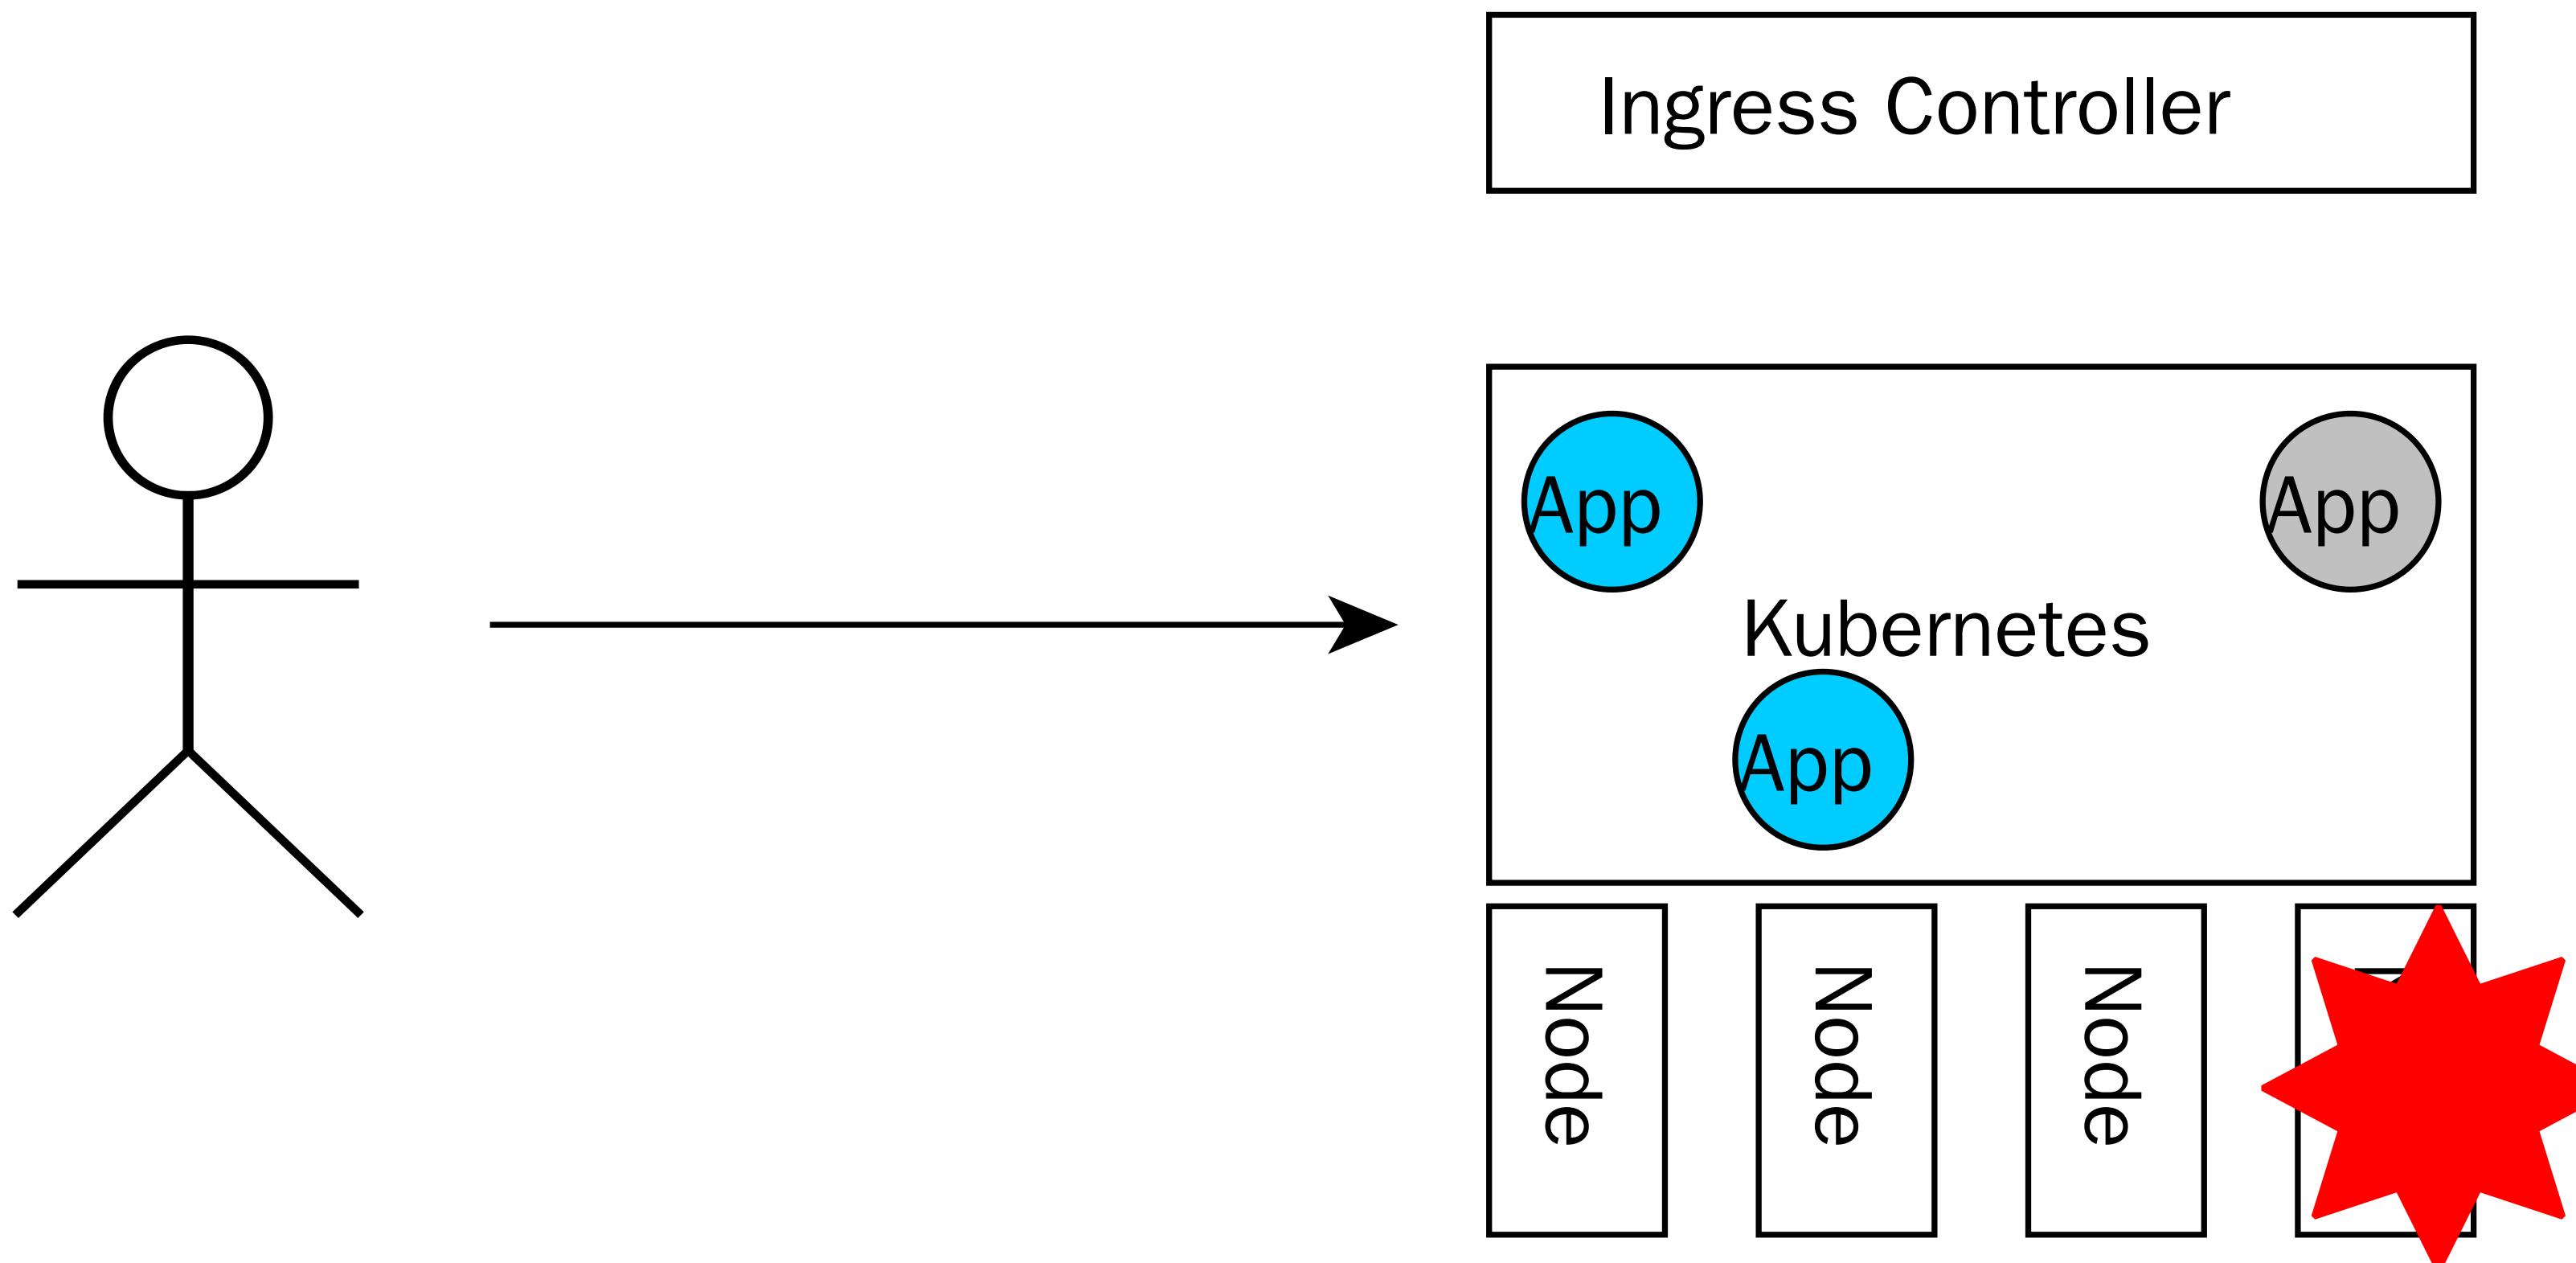
`make docker_push; kubectl create -f app-srv-dpl.yaml`

SCALE UP! SCALE DOWN!



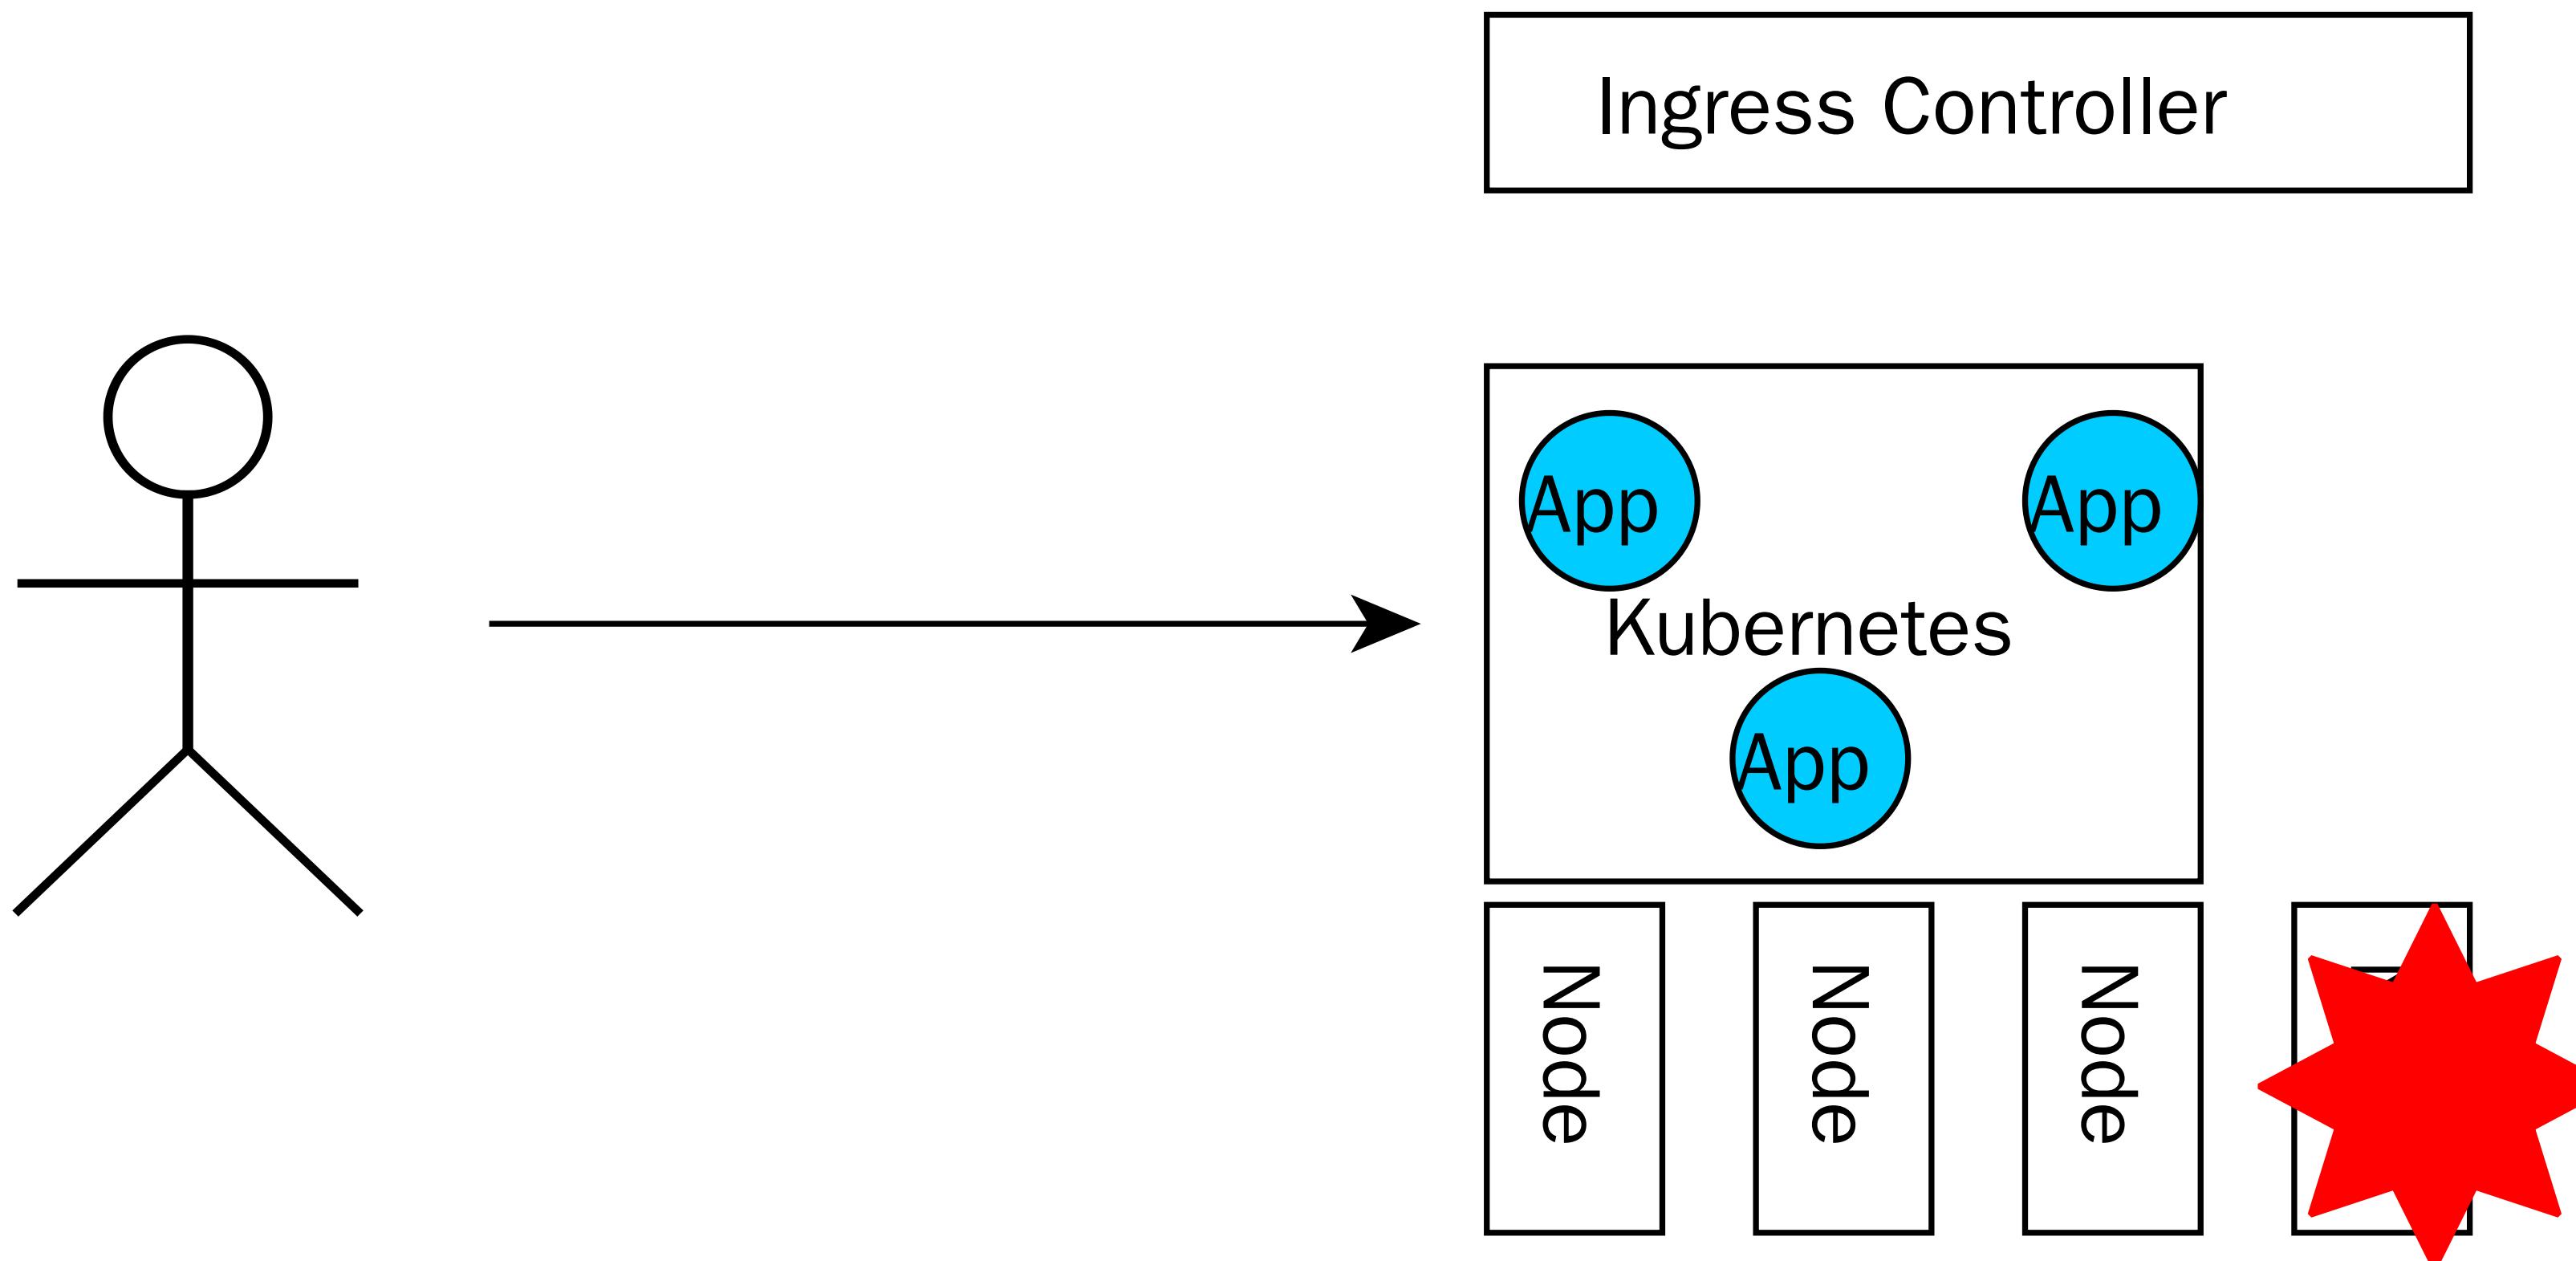
```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

RESILIENCE!



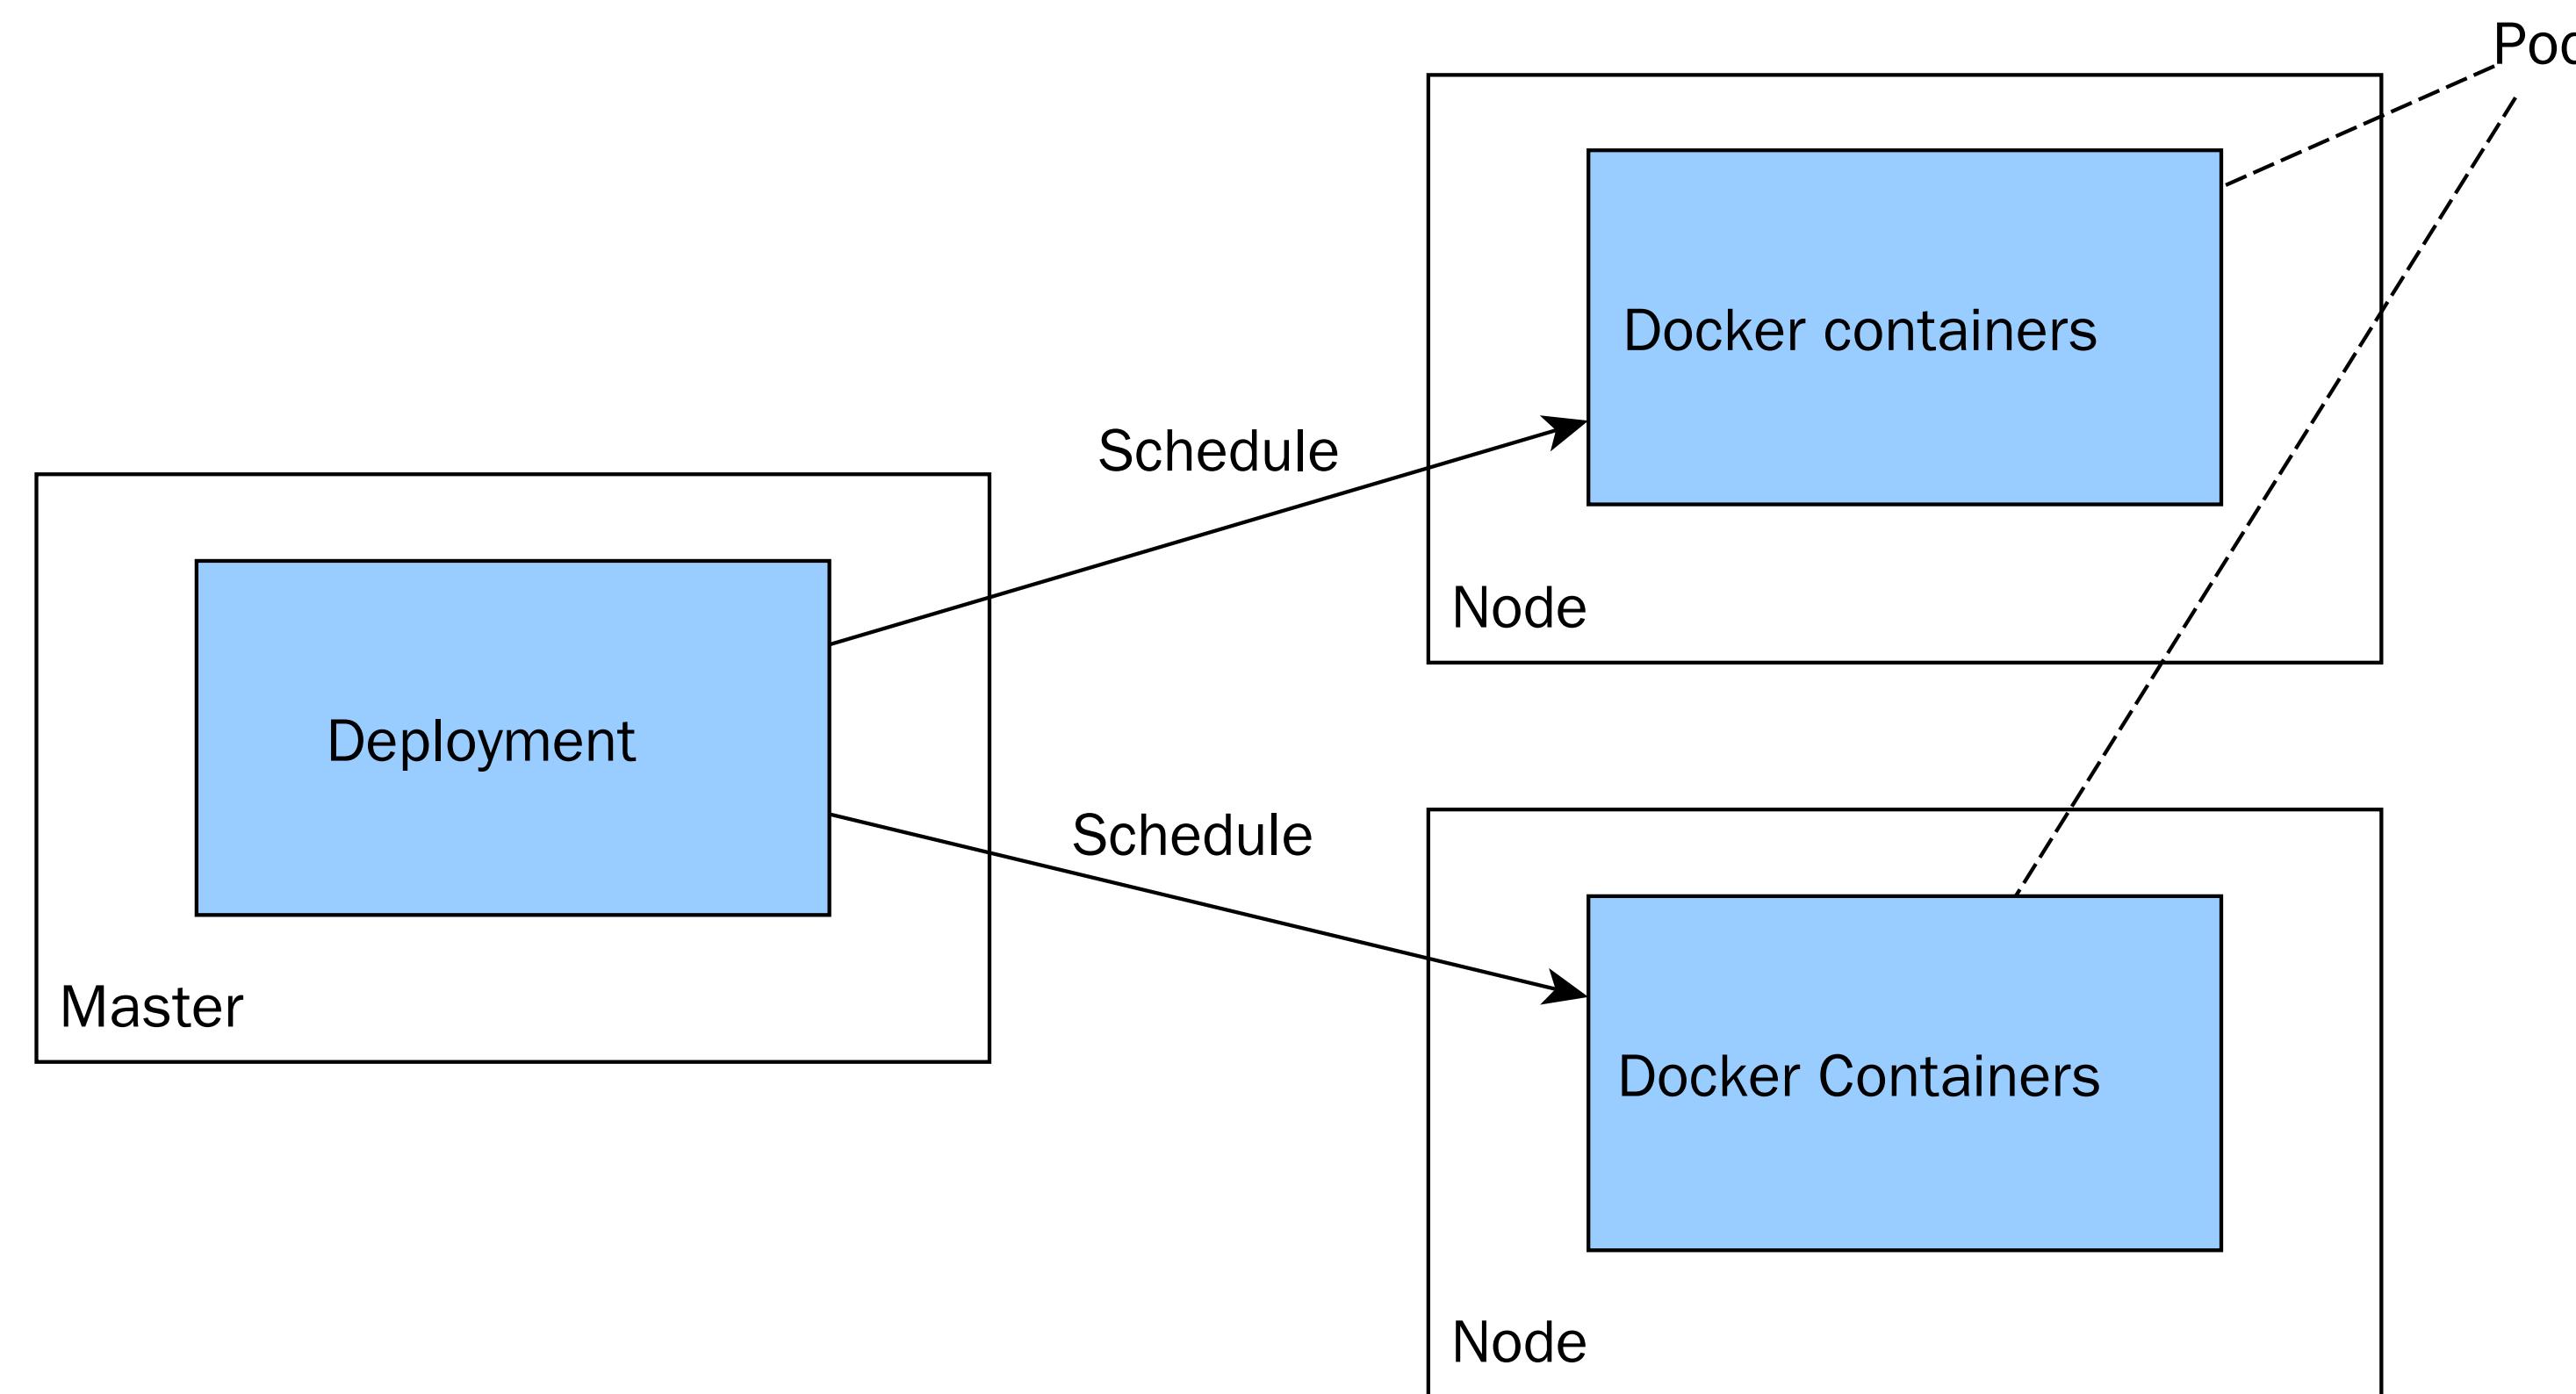
```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

RESILIENCE!



```
kubectl --replicas=3 -f app-srv-dpl.yaml
```

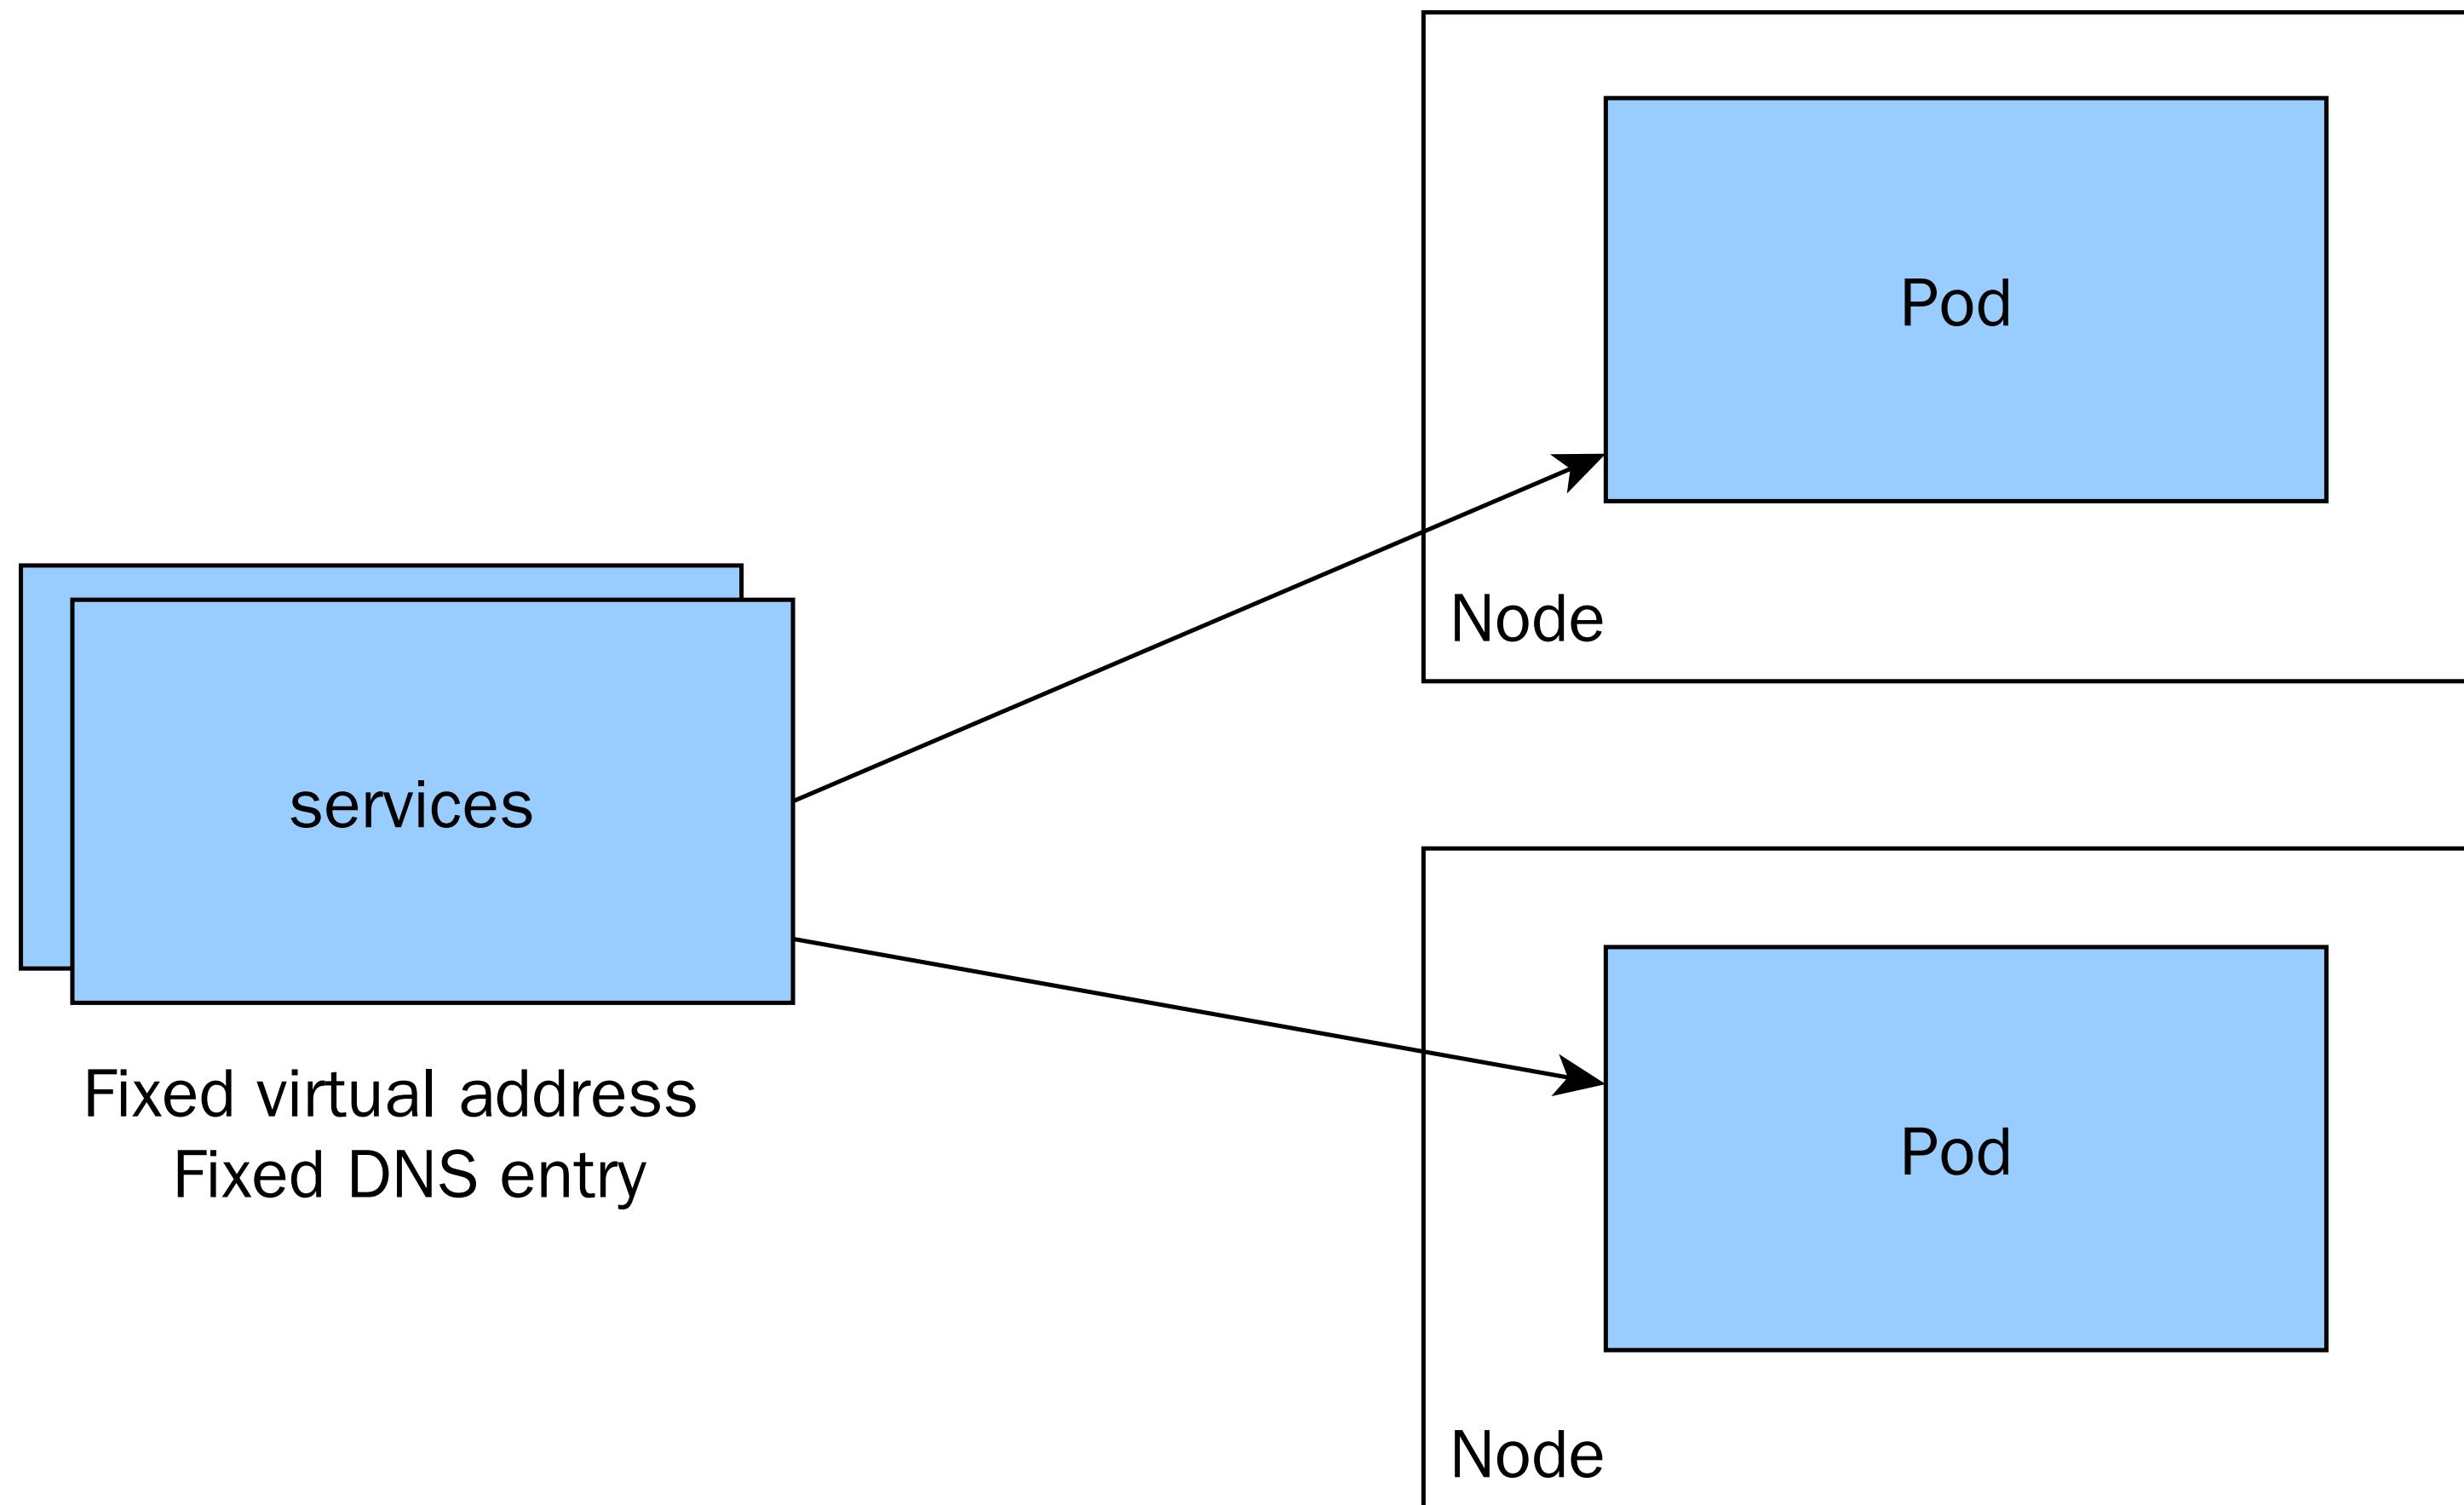
DEPLOYMENT AND PODS



deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-api
  labels:
    app: demo-api
spec:
  replicas: 3
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: demo-api
  template:
```

SERVICE AND PODS



Service matches pods based on labels

service.yml

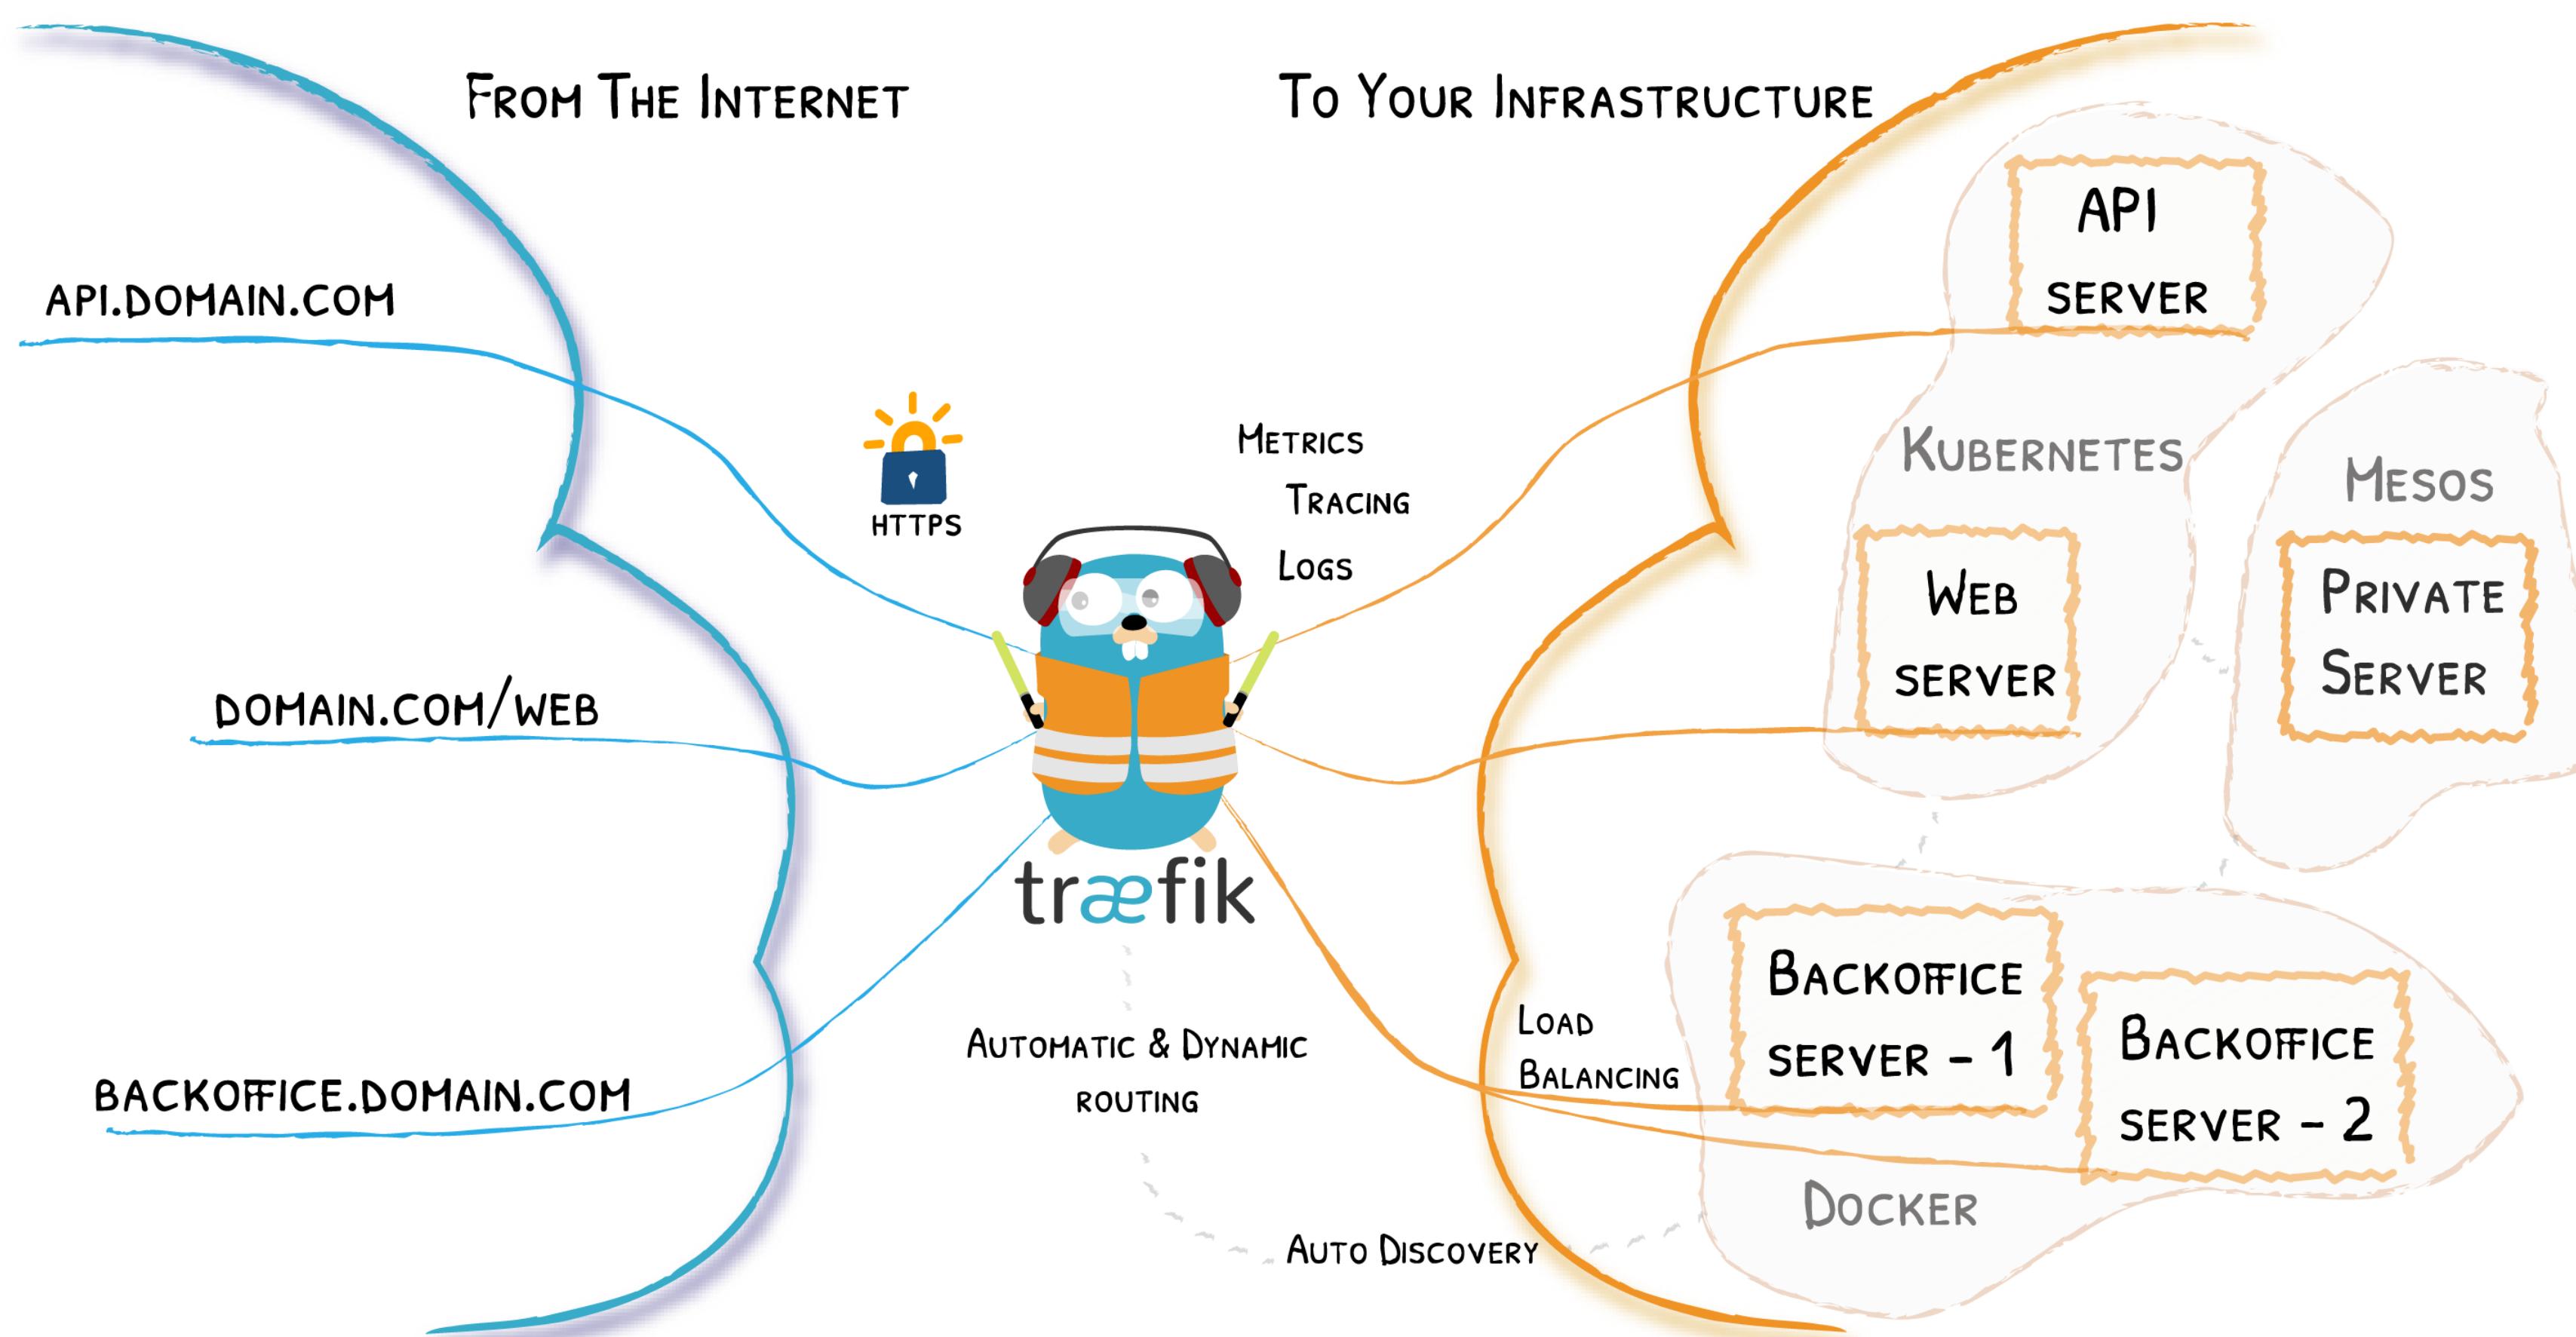
```
apiVersion: v1
kind: Service
metadata:
  name: demo-api
spec:
  ports:
  - port: 80
    protocol: TCP
  selector:
    app: demo-api
  type: LoadBalancer
```

SERVICE

```
curl http://users/list
```

Load balanced inside cluster!

HOW GET USER REQUESTS?



INGRESS

Pattern

api.smacc.io/v1/users users-v1

api.smacc.io/v2/users users-v2

smacc.io

web

BASIC CONCEPTS

Name	Purpose	
Service	Interface	Entry point (Service Name)
Deployment	Factory	How many pods, which pods
Pod	Implementation	1+ docker running

SERVICE DISCOVERY AND LABELING

- names in DNS:

curl http://users/list

- labels:

name=value

- annotations:

prometheus.io/scrape: "true"

SERVICE DISCOVERY

- auto-wiring with logging and monitoring
- drop-in installation (traefik, prometheus, ...)

INTEGRATION WITH MONITORING

```
apiVersion: v1
kind: Service
metadata:
  name: iam
  labels:
    app: iam
    tier: backend
  annotations:
    prometheus.io/scrape: "true"
    prometheus.io/scheme: "http"
    prometheus.io/path: "metrics"
    prometheus.io/port: "8088"
```

INGRESS CONFIGURATION

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: iam
  annotations:
    kubernetes.io/ingress.class: traefik
    traefik.frontend.rule.type: PathPrefixStrip
spec:
  rules:
  - host: example.org
    http:
      paths:
      - path: /iam
        backend:
```

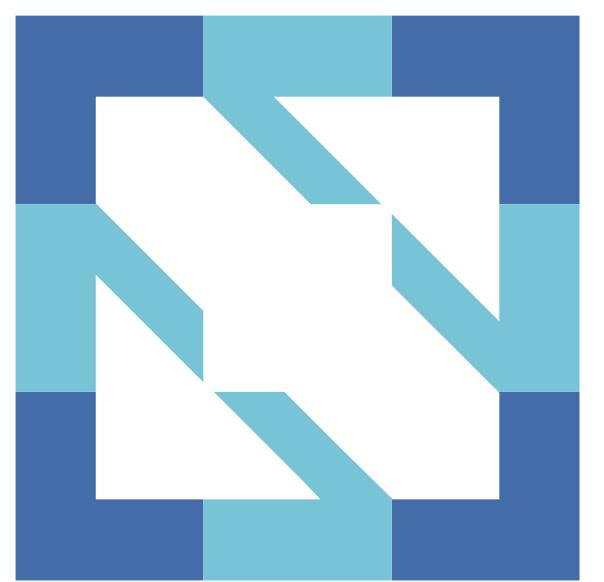
All hail the git!

- Yaml
- integration:
 - monitoring, alarming, ingress-controller
- ...
- Devs can forget about infra... almost
- DevOps Culture Dream!
 - + all tools -> CRD

simple semantic

- Complex release strategies are (relatively) easy
- Continuous Deployment is simply to implement!

github.com/wojciech12/talk_zero_downtime_deployment_with_kubernetes



**CLOUD NATIVE
COMPUTING FOUNDATION**

PROMETHEUS



Prometheus

Best in class monitoring and alerting

FLUENTD



fluentd

combined with ElasticSearch and Kibana

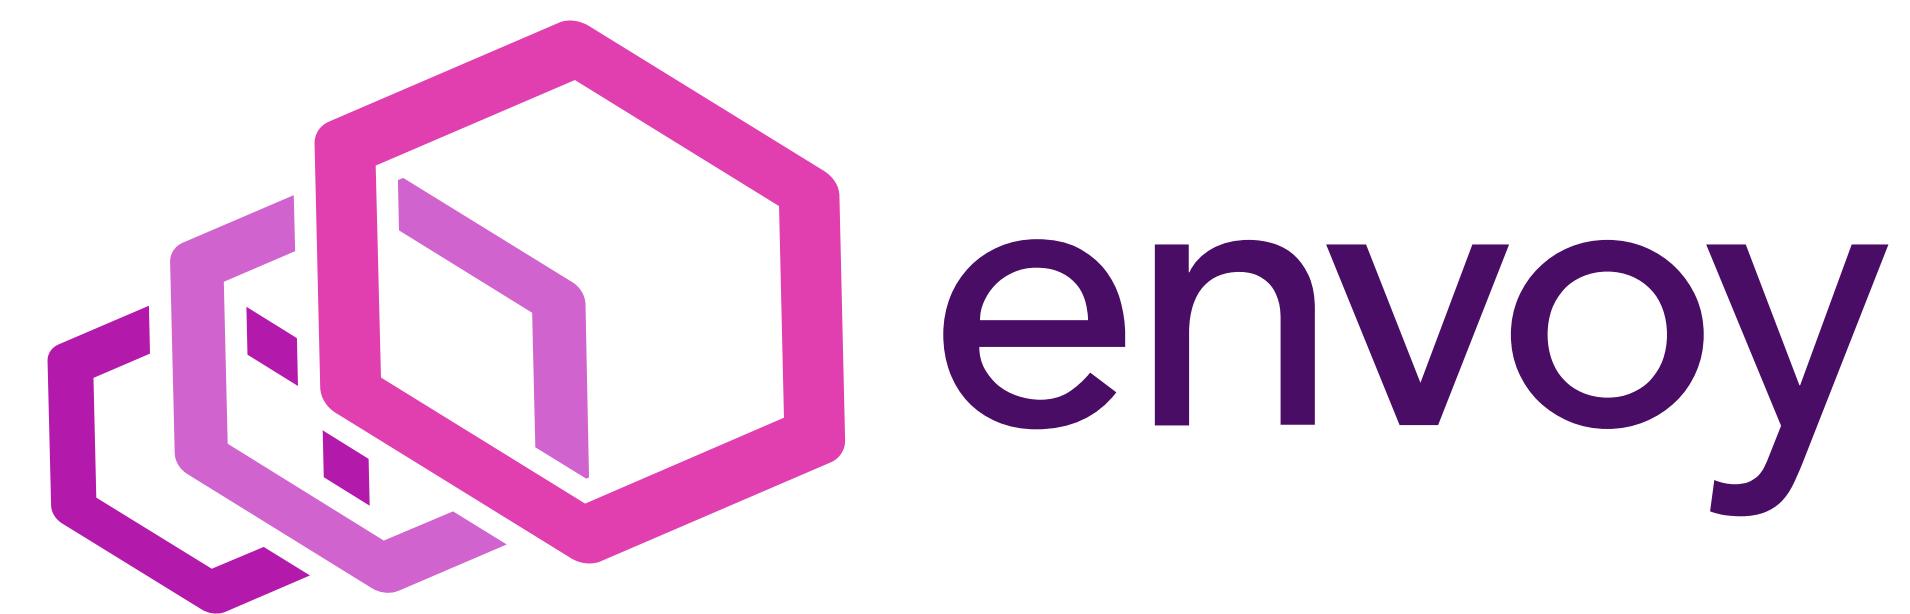
ETCD

Distributed key-value store.



Alternative: Consule, Zookeeper

INGRESS/PROXY



other: Nginx, Traefik

HOTTEST!



Not in CloudNative Foundation

TRACING



JAEGER

the last piece for observability, read [this blog](#)

K8S PACKAGE MANAGER



I do not recommend to use

MUCH MORE

- Kubeless/OpenFaaS for Serverless
- Kubeflow/Pachyderm for Machine Learning
- VirtualMachines with K8S - kata, kubevirt

COMMON

- Simple (reasonable) to run
- Focus on simplicity in operating at scale

HOW TO START

AS-A-SERVICE

- Google - GCE
- Azure - AKS
- Amazon - EKS

A lot of installers as well. See also [rancher](#).

AT HOME

- Mini-kube
- ...

A lot of installers as well



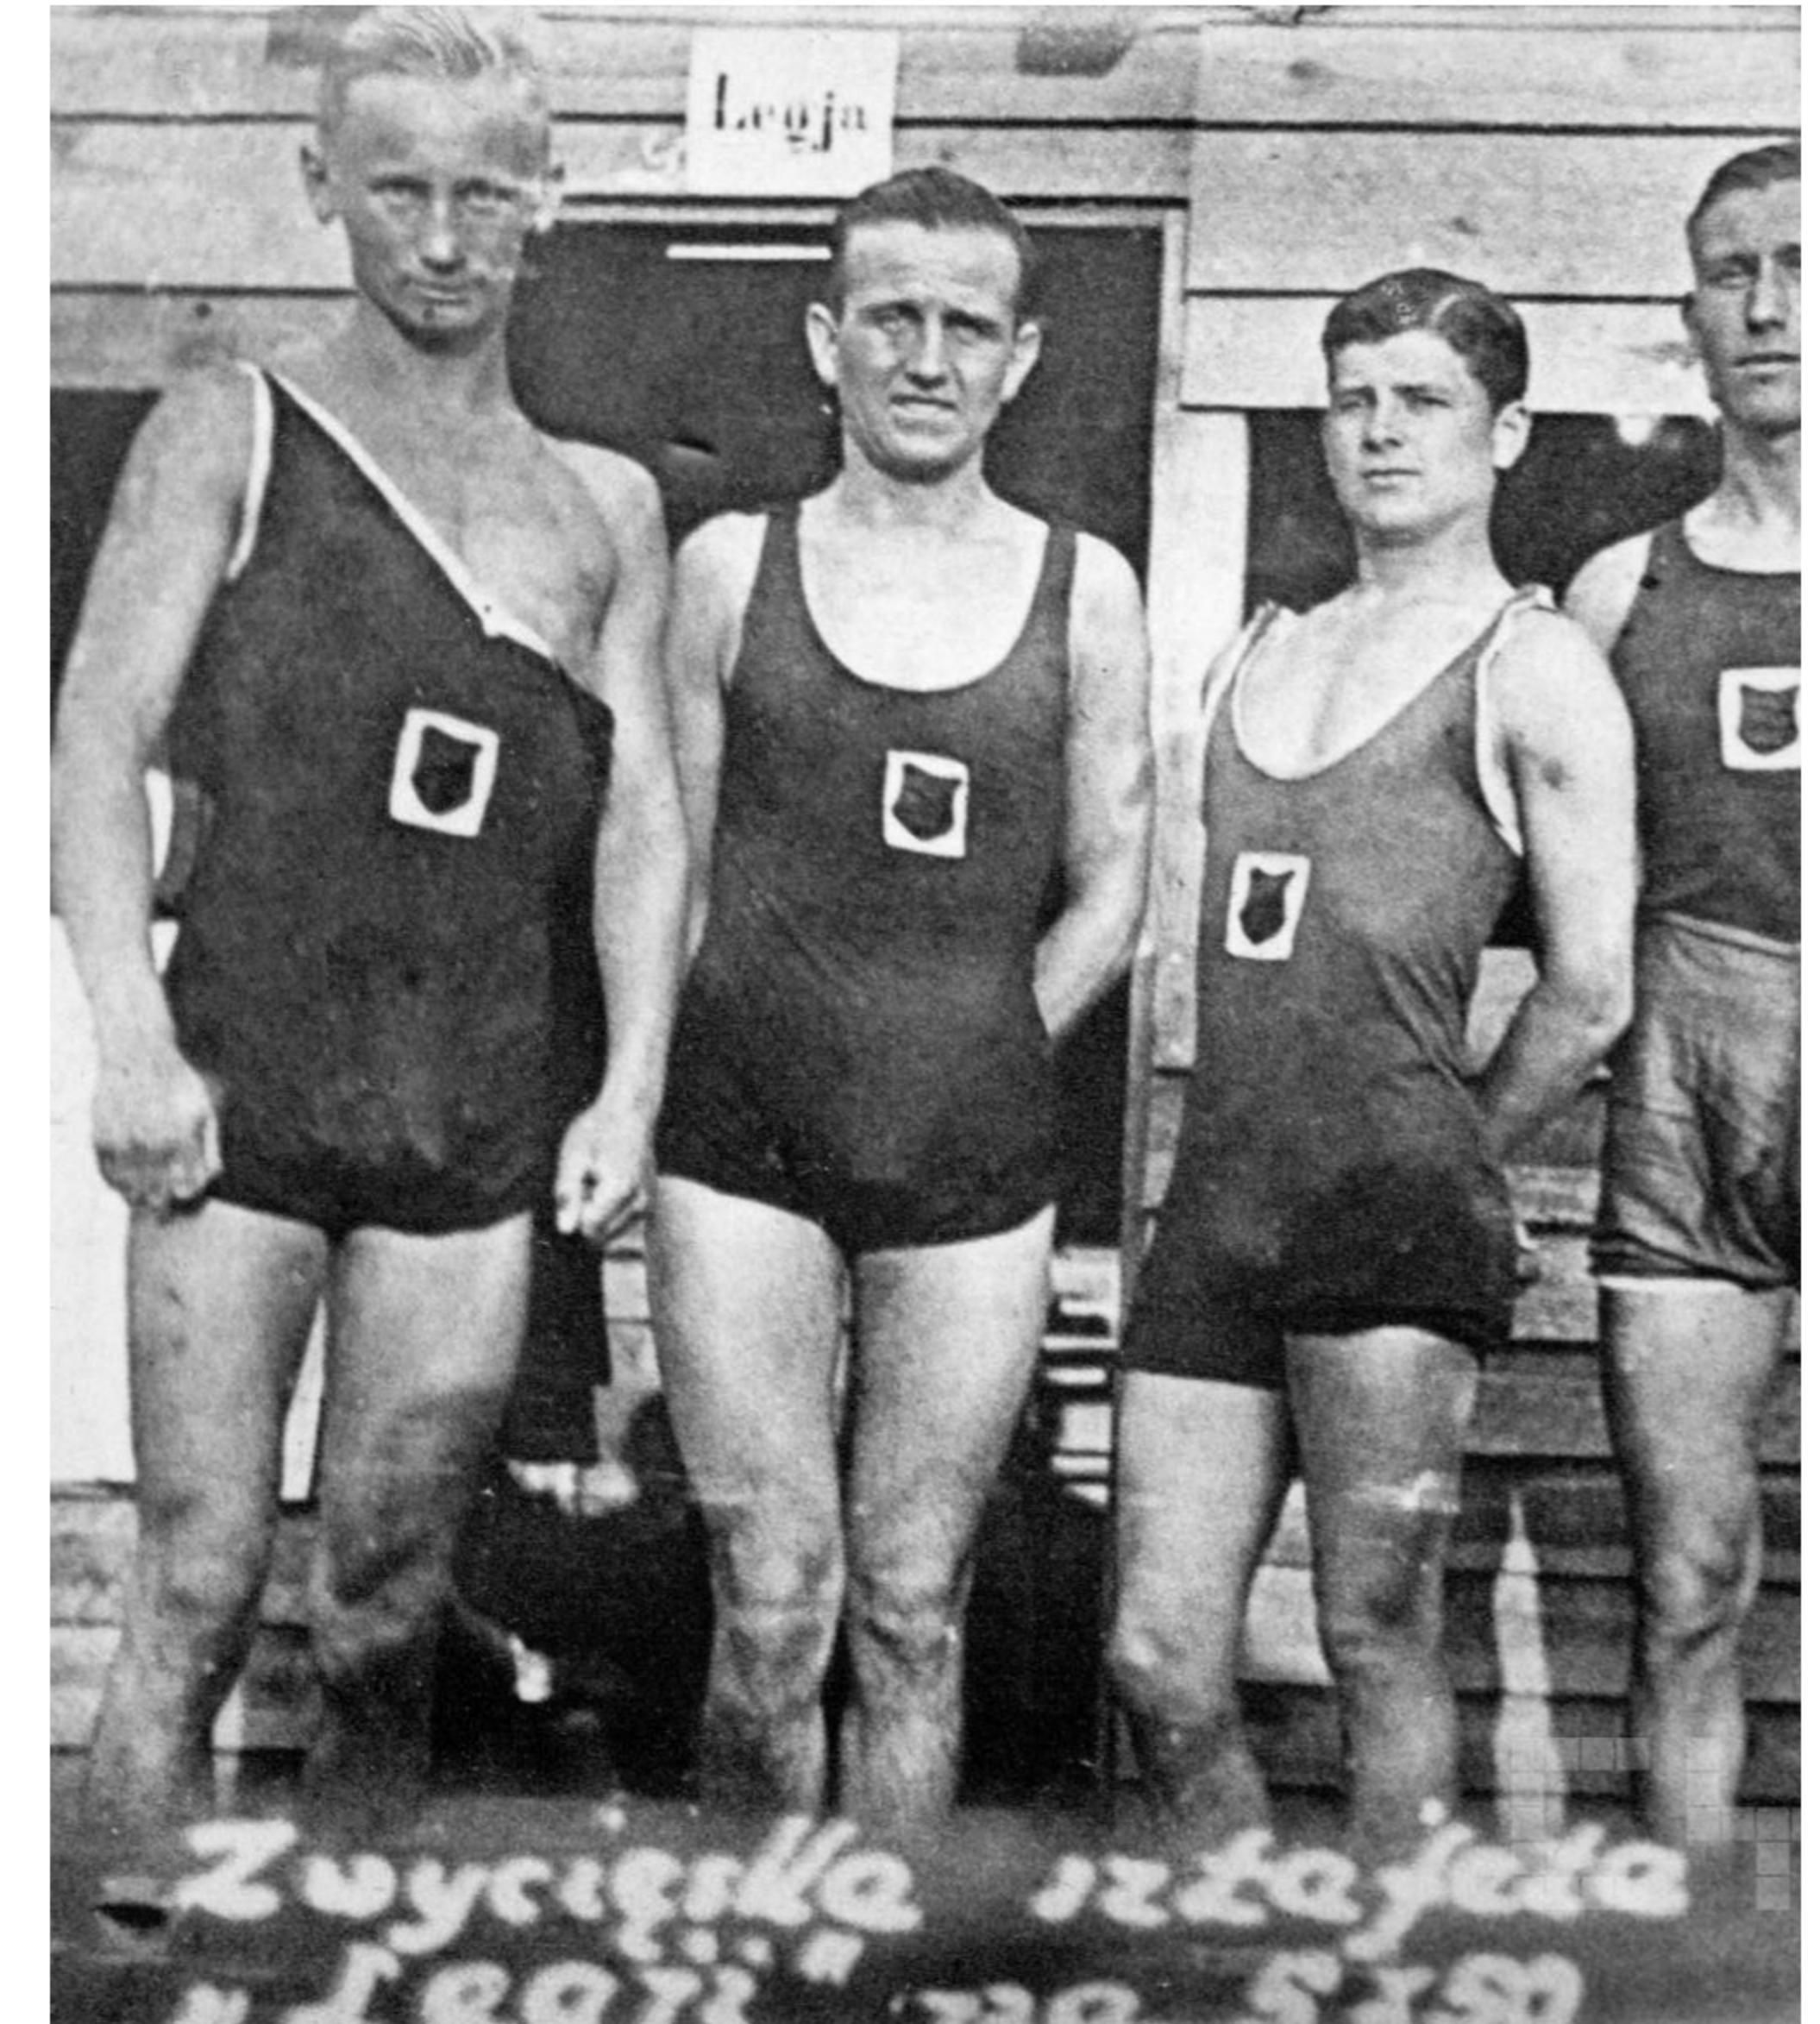
Keep it simple

BRUNO
FOTOGRAFIA

https://www.flickr.com/photos/bruno_brujah/

Keep everybody in the process

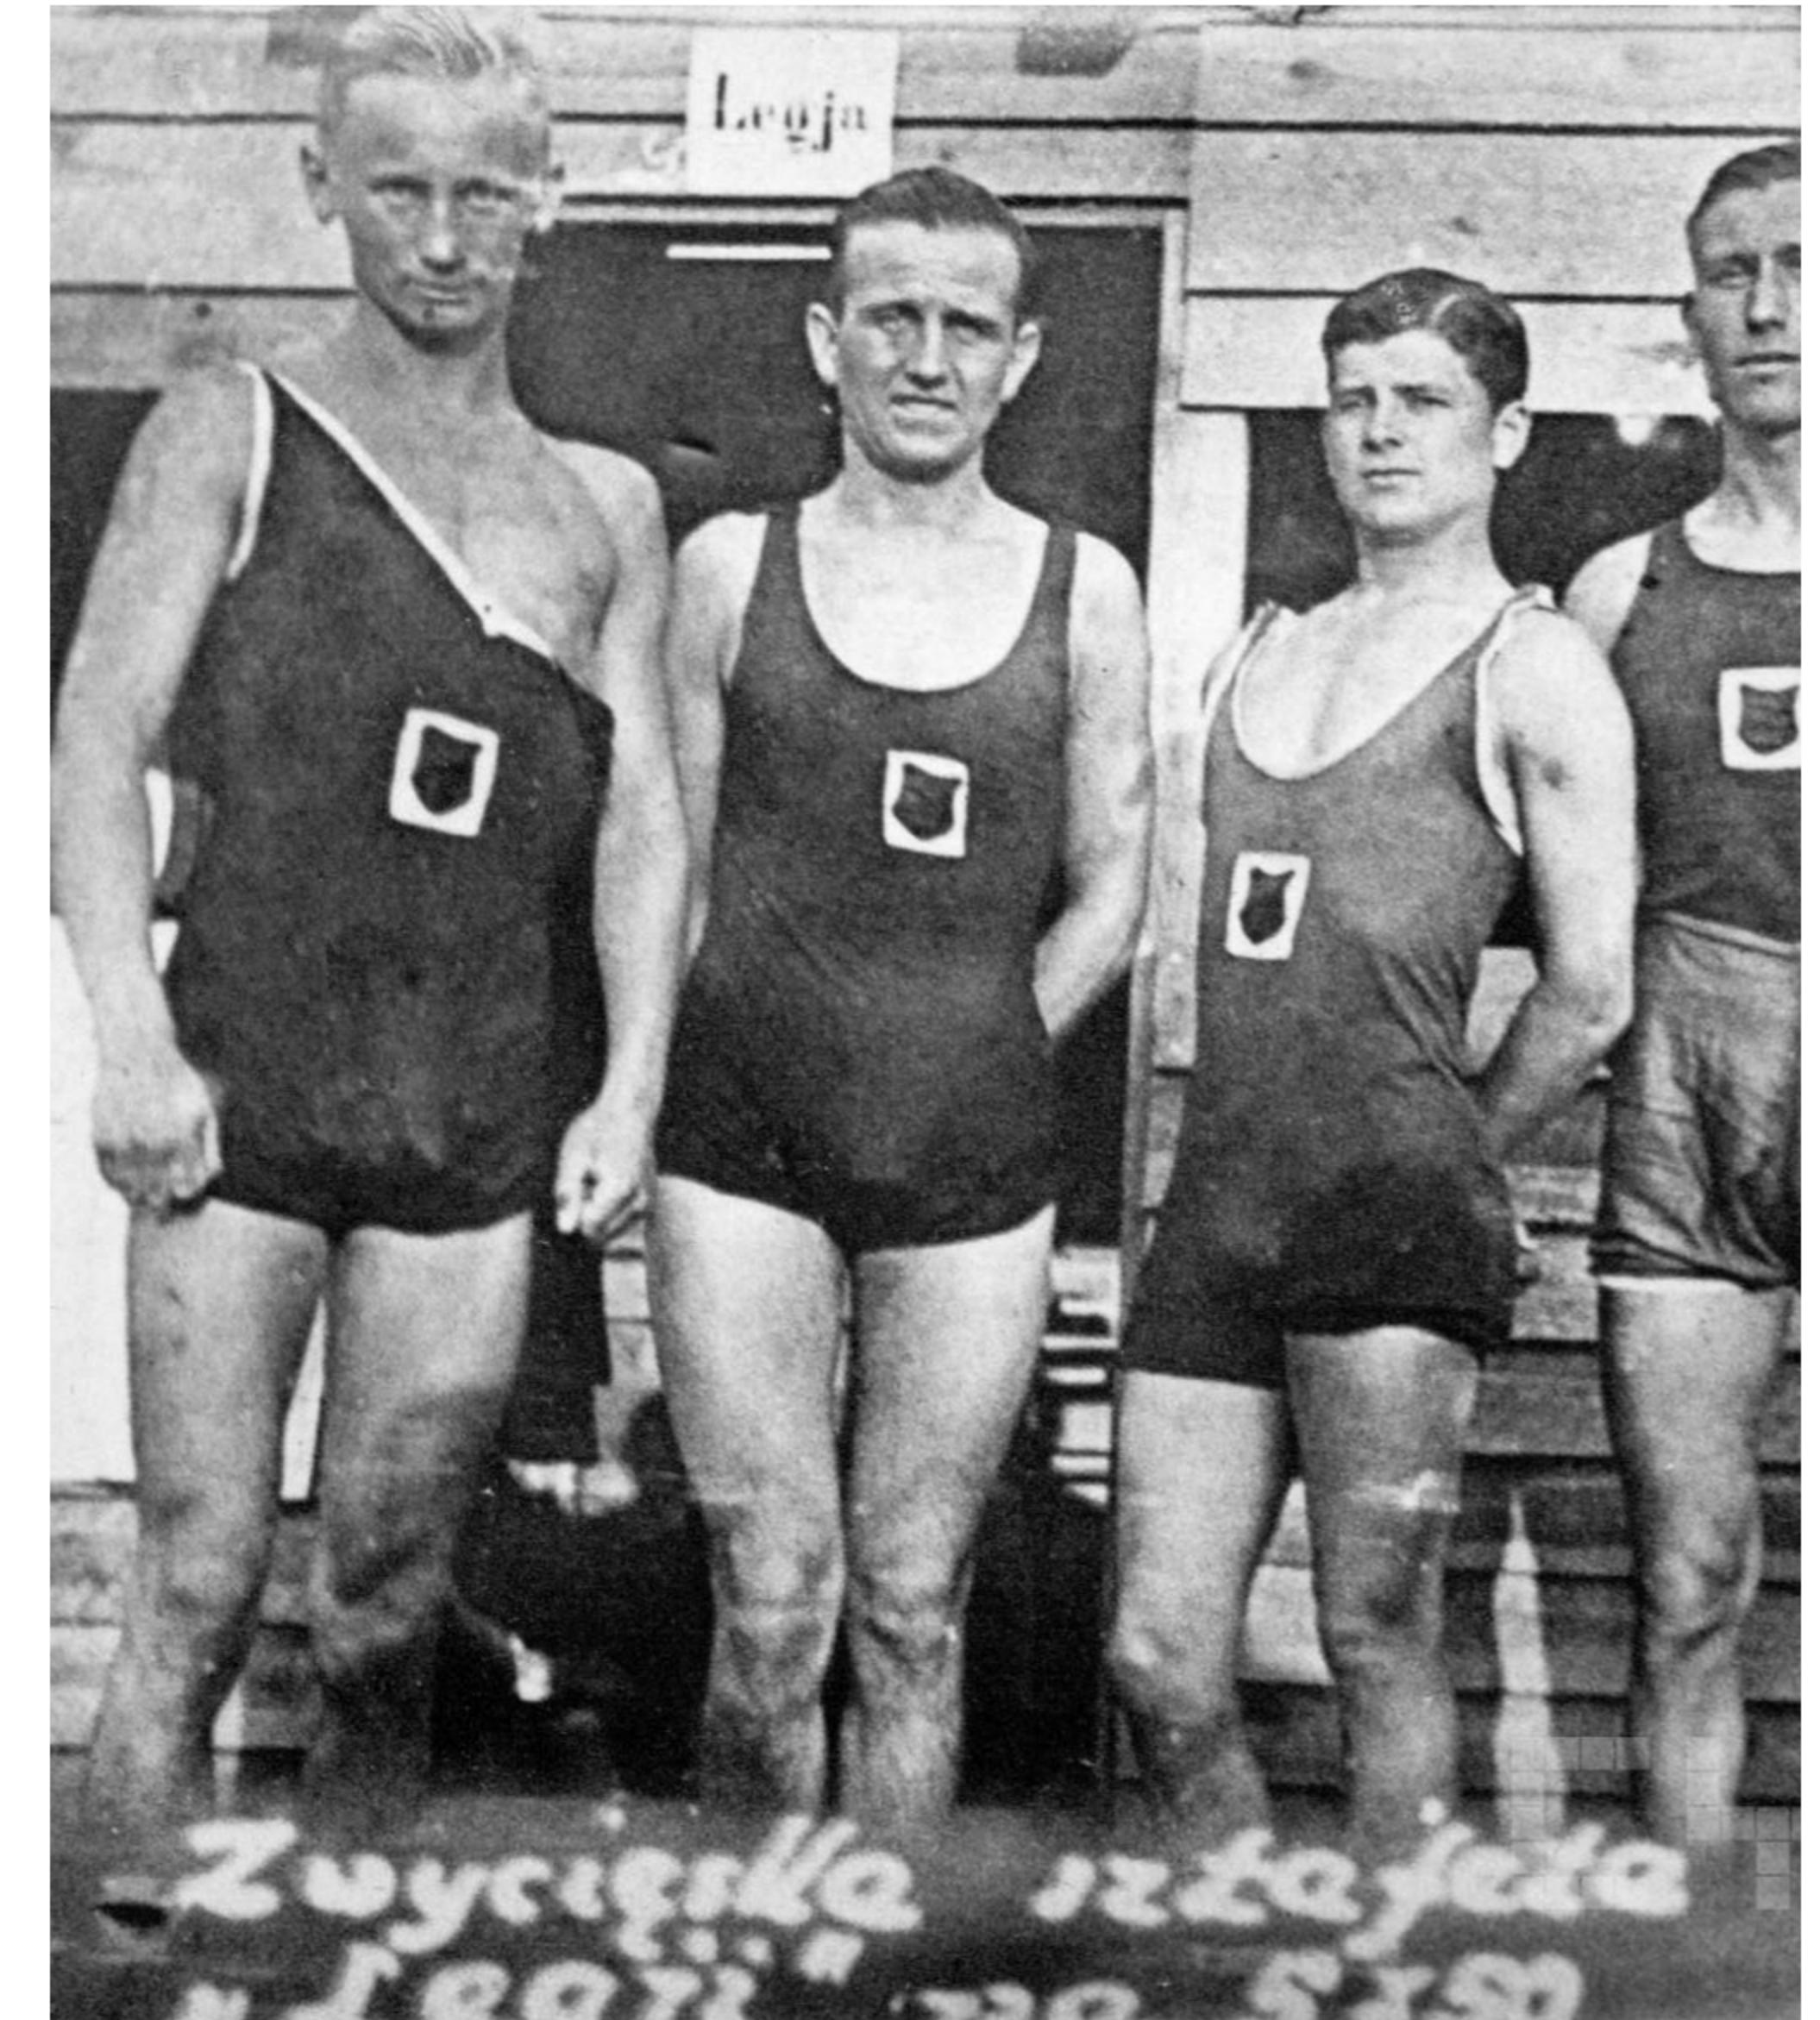
1. Teach the team
Kubernetes definitions
2. Keep the process
understandable



Keep everybody in the process

3. Keep K8S files, ...
easy to read

4. Do not terrorize with
how-amazing-
Kubernetes-is :D



and hide it!

your-platform.git

- `git tag` and `push`
- branches: dev, staging, master
- Generated k8s files
- PR to production

Similar to the [Kelsey Hightower approach](#)

KUBERNETES

- Pure and generated from simple templates with \${ENV_VAR}
- 2x kubernetes operators

GIT REPO

```
| - tools
|   | - kube-service.yaml
|   \- kube-deployment.yaml
|
| - Dockerfile
| - VERSION
\ - Makefile
```

Makefile

```
SERVICE_NAME=v-connector
DOCKER_REGISTRY=eu.gcr.io

test: test_short test_integration

run_local:

release: docker_build docker_push

kube_create_config:

deploy:
```

Copy&Paste from the project to project

Effective Kubernetes



https://www.flickr.com/photos/bruno_brujah/

- Start with small iterations
- Learn-as-you-go

Effective Kubernetes



https://www.flickr.com/photos/bruno_brujah/

- Simplicity
- Keep your Kubernetes understandable

*Hope k8s team keep it
this way*

Kubernetes VS Startup

- VM with application run in dockers
- After MVP version X, go to k8s
- Learn as-you-go
- Do not get high on all CN projects :)

THANK YOU. QUESTIONS?

btw. <https://github.com/hypatos/jobs>

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```

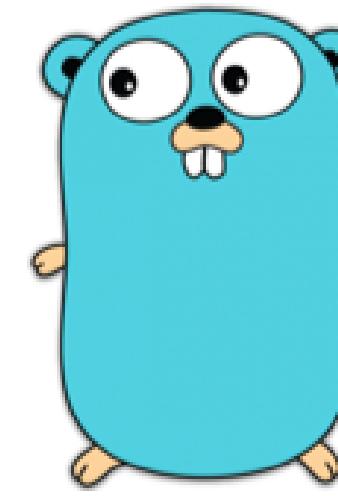


github.com/wojciech12

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



Hypatos



Go



BACKUP SLIDES

```
123 def distance_matrix(regions):  
124     """ Computes a distance matrix against a region list """  
125     tuples = [r.as_tuple() for r in regions]  
126     return cdist(tuples, tuples, region_distance)  
127  
128  
129 def clusterize(words, **kwargs):  
130     # TODO: write a cool docstring here  
131     db = DBSCAN(metric="precomputed", **kwargs)  
132     X = distance_matrix([Region.from_word(w) for w in words])  
133     labels = [int(l) for l in db.fit_predict(X)]
```



INFRA TOOLS

- Prometheus + AlertMnager + Grafana
- Traefik
- Kafka - Yolean/kubernetes-kafka
- Vault on Etcd - banzaicloud/bank-vaults

DATA STORES

- Kafka - Yolean/kubernetes-kafka
- Etcd - coreos/etcd-operator
- DB: PSQL and Mongo

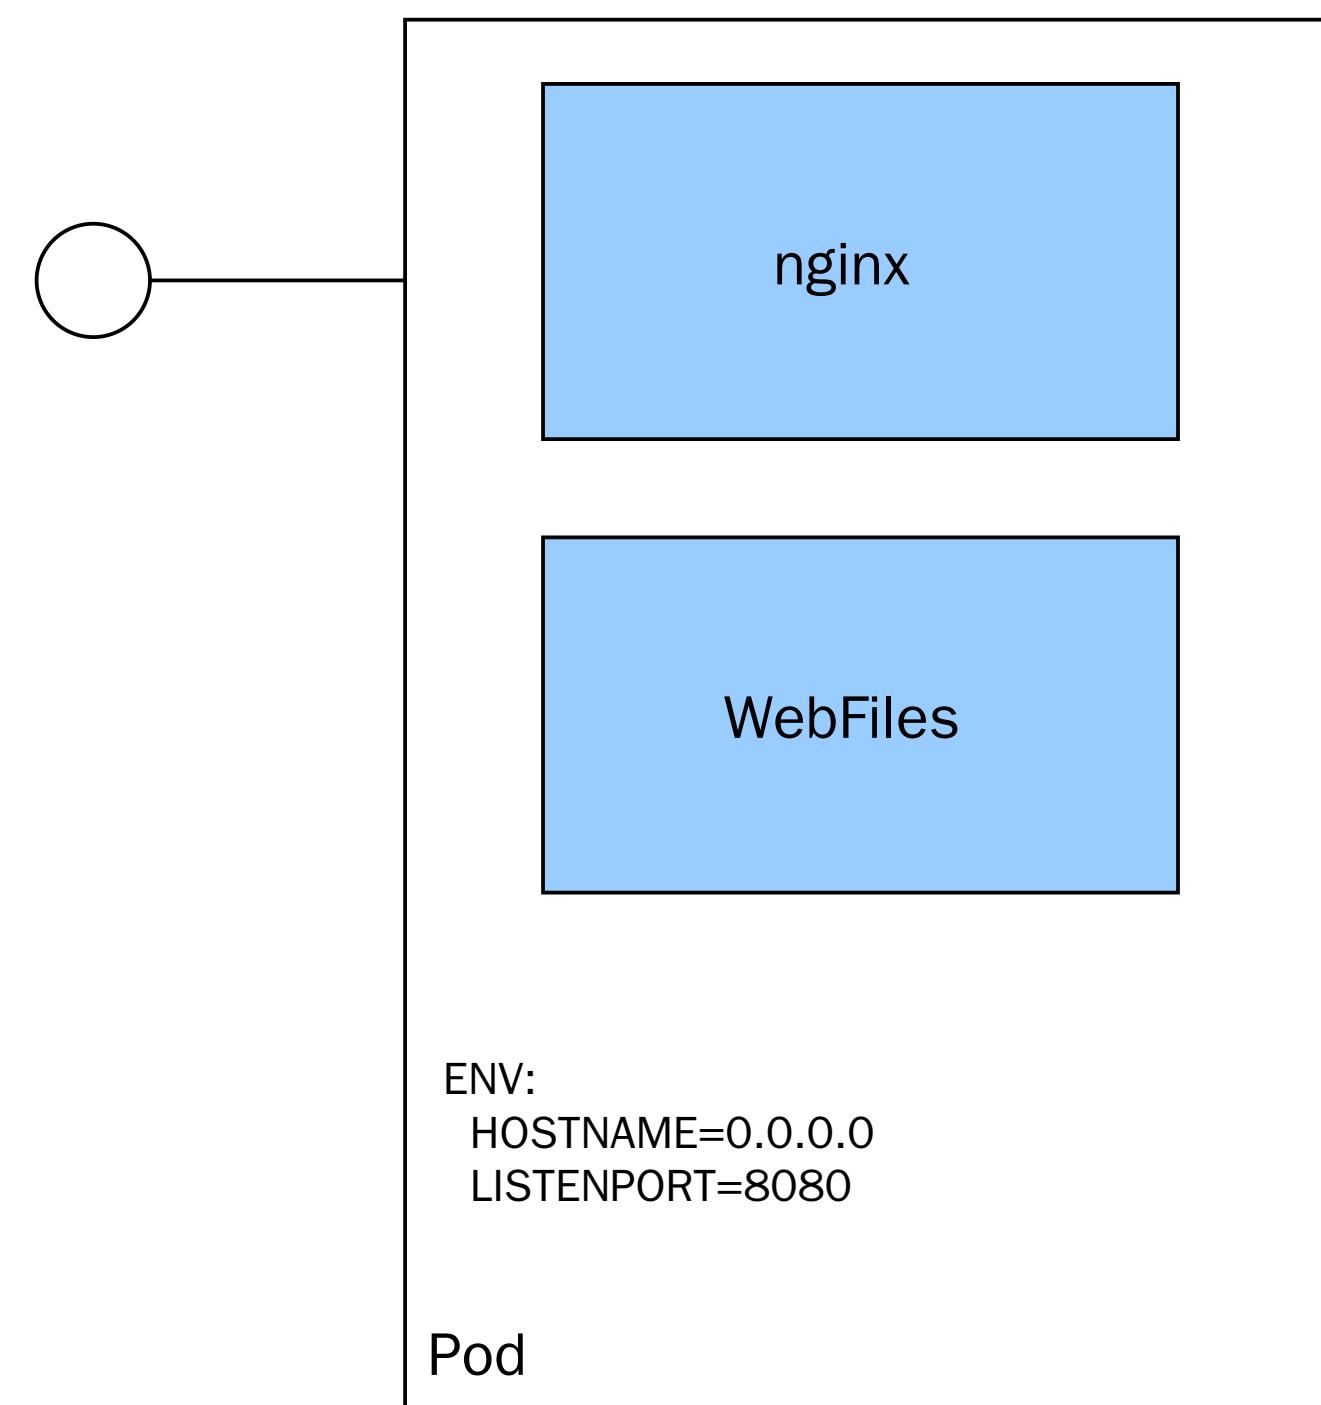
BACKUPS:

- old-school backups
- replications

KUBERNETES CONCEPTS+

PODS

- See each other on localhost
- Live and die together
- Can expose multiple ports



SIDE-CARS

